

homework6-skel

October 24, 2020

NAME: Nigel Mansell

SECTION: 995

CS 5970: Machine Learning Practices

1 Homework 6: Cross Validation

1.1 Assignment Overview

First read through the entire notebook, do not write any code. This assignment is more complex than previous, and it will be helpful to have a sense of the structure before you start coding.

Follow the TODOs and read through and understand any provided code.

All the plotting functions have been provided. You should not need to alter any of these.

1.1.1 Task

For this assignment you will be implementing **holistic cross validation**. Cross validation is a procedure that involves training, validating, and testing a model on different subsets of the data set to evaluate how well the model will generalize to unseen examples. Additionally, cross validation is a good tool for evaluating models when only small amounts of data are available.

The train sets are utilized for the various models to learn with, the validation sets are utilized to initially evaluate and select the best performing model. The test sets are utilized to determine how well the **chosen model** actually will generalize to unseen examples.

The validation and test sets can often seem similar conceptually, however, the key difference is that the validation performance is used to actually make guided decisions about model tuning (i.e., hyper-parameter values). Decisions about which hyper-parameters to use are never done based on the test set. The test set performance evaluates the generalized performance on data unused for hyper-parameter selection and training.

1.1.2 Data set

The BMI data will be utilized. Recall: * *MI* files contain data with the number of activations for 48 neurons, at multiple time points, for a single fold. There are 20 folds (20 files), where each fold consists of over 1000 time points (the rows). At each time point, we record the number of activations for each neuron for 20 bins. Therefore, each time point has $48 * 20 = 960$ columns.

* *theta* files record the angular position of the shoulder (in column 0) and the elbow (in column 1) for each time point.

* *dtheta* files record the angular velocity of the shoulder (in column 0) and the elbow (in column

1) for each time point.

* *torque* files record the torque of the shoulder (in column 0) and the elbow (in column 1) for each time point.

* *time* files record the actual time stamp of each time point.

1.1.3 Objectives

- Implement and understand **holistic cross validation**
- Training set size sensitivity analysis

1.1.4 Notes

- Do not save work within the ml_practices folder

1.1.5 General References

- [Guide to Jupyter](#)
- [Python Built-in Functions](#)
- [Python Data Structures](#)
- [Numpy Reference](#)
- [Numpy Cheat Sheet](#)
- [Summary of matplotlib](#)
- [DataCamp: Matplotlib](#)
- [Pandas DataFrames](#)
- [Sci-kit Learn Linear Models](#)
- [Sci-kit Learn Ensemble Models](#)
- [Sci-kit Learn Metrics](#)
- [Sci-kit Learn Model Selection](#)

1.1.6 Hand-In Procedure

- Execute all cells so they are showing correct results
- Notebook:
 - Submit this file (.ipynb) to the Canvas HW6 dropbox
- PDF:
 - File/Print/Print to file -> Produces a copy of the notebook in PDF format
 - Submit the PDF file to the Gradescope HW6 dropbox

```
[2]: import pandas as pd
import numpy as np
import scipy.stats as stats
import os, re, fnmatch
import pathlib, itertools, time
import matplotlib.pyplot as plt
import joblib

from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.metrics import explained_variance_score
from sklearn.linear_model import ElasticNet
```

```

FIGW = 10
FIGH = 6
FONTSIZE = 12

HOME_DIR = pathlib.Path.home()

plt.rcParams['figure.figsize'] = (FIGW, FIGH)
plt.rcParams['font.size'] = FONTSIZE

plt.rcParams['xtick.labelsize'] = FONTSIZE
plt.rcParams['ytick.labelsize'] = FONTSIZE

%matplotlib inline

```

```

[3]: """
    Display current working directory of this notebook. If you are using relative
    paths for your data, then it needs to be relative to the CWD.
    """
    pathlib.Path.cwd()

```

```

[3]: PosixPath('/home/nigel/Desktop/mlp/homework6')

```

2 LOAD DATA

```

[4]: def read_bmi_file_set(directory, filebase):
    """
    Read a set of CSV files and append them together
    :param directory: The directory in which to scan for the CSV files
    :param filebase: A file specification that potentially includes wildcards
    :returns: A list of Numpy arrays (one for each fold)
    """

    # The set of files in the directory
    files = fnmatch.filter(os.listdir(directory), filebase)
    files.sort()

    # Create a list of Pandas objects; each from a file in the directory that
    ↳ matches filebase
    lst = [pd.read_csv(directory + "/" + file, delim_whitespace=True,
    ↳ header=None).values for file in files]

    # Concatenate the Pandas objects together. ignore_index is critical here,
    ↳ so that
    # the duplicate row indices are addressed

```

```
return lst
```

```
[5]: """ PROVIDED
Load the BMI data from all the folds, using read_bmi_file_set()
"""
# TODO: might need to change; assumes ml_practices is in home directory
#changed to my directory
dir_name = str(HOME_DIR / '/home/nigel/Desktop/mlp/mlp_2020/datasets/bmi/
↳DAT6_08')

MI_folds = read_bmi_file_set(dir_name, 'MI_fold*')
theta_folds = read_bmi_file_set(dir_name, 'theta_fold*')
dtheta_folds = read_bmi_file_set(dir_name, 'dtheta_fold*')
torque_folds = read_bmi_file_set(dir_name, 'torque_fold*')
time_folds = read_bmi_file_set(dir_name, 'time_fold*')

alldata_folds = zip(MI_folds, theta_folds, dtheta_folds,
                    torque_folds, time_folds)

nfolders = len(MI_folds)
nfolders
```

[5]: 20

```
[6]: """ PROVIDED
Print out the shape of all the data for each fold
"""
for i, (MI, theta, dtheta, torque, time) in enumerate(alldata_folds):
    print("FOLD %2d " % i, MI.shape, theta.shape,
          dtheta.shape, torque.shape, time.shape)
```

```
FOLD 0 (1194, 960) (1194, 2) (1194, 2) (1194, 2) (1194, 1)
FOLD 1 (1105, 960) (1105, 2) (1105, 2) (1105, 2) (1105, 1)
FOLD 2 (1532, 960) (1532, 2) (1532, 2) (1532, 2) (1532, 1)
FOLD 3 (1266, 960) (1266, 2) (1266, 2) (1266, 2) (1266, 1)
FOLD 4 (1499, 960) (1499, 2) (1499, 2) (1499, 2) (1499, 1)
FOLD 5 (1253, 960) (1253, 2) (1253, 2) (1253, 2) (1253, 1)
FOLD 6 (1376, 960) (1376, 2) (1376, 2) (1376, 2) (1376, 1)
FOLD 7 (1131, 960) (1131, 2) (1131, 2) (1131, 2) (1131, 1)
FOLD 8 (1248, 960) (1248, 2) (1248, 2) (1248, 2) (1248, 1)
FOLD 9 (1258, 960) (1258, 2) (1258, 2) (1258, 2) (1258, 1)
FOLD 10 (1266, 960) (1266, 2) (1266, 2) (1266, 2) (1266, 1)
FOLD 11 (1147, 960) (1147, 2) (1147, 2) (1147, 2) (1147, 1)
FOLD 12 (1226, 960) (1226, 2) (1226, 2) (1226, 2) (1226, 1)
FOLD 13 (1239, 960) (1239, 2) (1239, 2) (1239, 2) (1239, 1)
FOLD 14 (1571, 960) (1571, 2) (1571, 2) (1571, 2) (1571, 1)
FOLD 15 (1360, 960) (1360, 2) (1360, 2) (1360, 2) (1360, 1)
```

```
FOLD 16 (1580, 960) (1580, 2) (1580, 2) (1580, 2) (1580, 1)
FOLD 17 (1365, 960) (1365, 2) (1365, 2) (1365, 2) (1365, 1)
FOLD 18 (1390, 960) (1390, 2) (1390, 2) (1390, 2) (1390, 1)
FOLD 19 (1290, 960) (1290, 2) (1290, 2) (1290, 2) (1290, 1)
```

3 PARAMETER SET LIST

```
[7]: """ PROVIDED
Construct the Cartesian product of the parameters
"""
def generate_paramsets(param_lists):
    """
    Construct the Cartesian product of the parameters
    PARAMS:
        params_lists: dict of lists of values to try for each parameter.
                      keys of the dict are the names of the parameters
                      values are lists of values to try for the
                      corresponding parameter
    RETURNS: a list of dicts that make up the Cartesian product of the
              parameters
    """
    keys, values = zip(*param_lists.items())
    # Determines cartesian product of parameter values
    combos = itertools.product(*values)
    # Constructs list of dictionaries
    combos_dicts = [dict(zip(keys, vals)) for vals in combos]
    return list(combos_dicts)
```

4 PERFORMANCE EVALUATION

```
[8]: """ PROVIDED
Evaluate the performance of an already trained model on some data
"""
def mse_rmse(trues, preds):
    """
    Compute MSE and rMSE for each column separately.
    """
    mse = np.sum(np.square(trues - preds), axis=0) / trues.shape[0]
    rmse = np.sqrt(mse)
    return mse, rmse

""" TODO
Finish implementation by just returning the dictionary of results
"""
def score_eval(model, X, y, preds):
    """
```

Compute the model predictions and corresponding scores, for an already trained model.

PARAMS:

model: model to predict with

X: input feature data

y: true output for X

preds: predicted output for X

RETURNS: results as a dictionary of numpy arrays

mse: mean squared error for each column

rmse: rMSE for each column

evan: explained variance, best is 1.0

score: score computed by the models score() method

'''

`score = model.score(X, y)`

`mse, rmse = mse_rmse(y, preds)`

`evan = explained_variance_score(y, preds)`

TODO: Complete the results dictionary. This is a

dictionary of numpy arrays. The numpy arrays must

be row vectors, where each element is the result

for a different output, when using multiple regression.

The keys of the dictionary are the name of the performance

metric, and the values are the numpy row vectors

#finished results dictionary

`results = {'mse': np.reshape(mse, (1, -1)),`

`'rmse': np.reshape(rmse, (1, -1)),`

`'evan': np.reshape(evan, (1, -1)),`

`'score': np.reshape(score, (1, -1))`

`}`

`return results`

5 CROSS VALIDATION

```
[9]: """ TODO
Complete KFoldHolisticCrossValidation implementation
General Procedure:
grid_cross_validation():
    for each hyper-parameter combination:
        set hyper-parameters of the model
        for each training set size:
            perform_cross_validation()
            record results for the hyper-parameter combination

perform_cross_validation():
```

```

    for each rotation:
        split data into train, test, val sets using get_data()
        train the model
        evaluate the model on train, val, and test sets
        record the results
    record results by size
"""
class KFoldHolisticCrossValidation():
    def __init__(self, model, paramsets, eval_func, opt_metric,
                  maximize_opt_metric=False, trainsizes=[1], rotation_skip=1):
        ''' TODO
        Object for managing and performing cross validation for a given model,
        → for
        a list of parameter sets and train set sizes. Note, train set size is,
        → in
        terms of number of folds (not samples)
        PARAMS:
            model: base ML model

            paramsets: list of dicts of parameter sets to give to the model

            eval_func: handle to function used to evaluate/score the model
                The eval_func must have the following arguments: model,
                X, ytrue, ypreds and return a dict of numpy arrays with
                shape 1-by-n, where n is the number of outputs if using
                multiple regression.
                template function header: eval_func(model, X, y, preds)
                template output: {'metrics1':1_by_n_array, ...}

            opt_metric: the optized metric. one of the metric key names
                returned from eval_func to use to pick the best
                parameter sets

            maximize_opt_metric: True if opt_metric is maximized; False if,
        → minimized

            trainsizes: list of training set sizes (in number of folds) to try

            rotation_skip: build model and evaluate every ith rotation (1=all
                possible rotations; 2=every other rotation, etc.)

        '''
        # TODO: set the class variables

        #setting the class variables
        self.model = model
        self.paramsets = paramsets

```

```

self.trainsizes = trainsizes
self.eval_func = eval_func
self.opt_metric = opt_metric + '_mean'
self.maximize_opt_metric = maximize_opt_metric
self.rotation_skip = rotation_skip

# Results attributes
# Full recording of all results for all paramsets, sizes, rotations,
# and metrics. This is a list of dictionaries for each paramset
self.results = None
# Validation summary report of all means and standard deviations for
# all metrics, for all paramsets, and sizes. This is a 3D s-by-r-by-p
# numpy array. Where s is the number of sizes, r the number of summary
# metrics +2, and p is the number of paramsets
self.report_by_size = None
# List of the indices of the best paramset for each size
self.best_param_inds = None

def get_data(self, all_Xfolds, all_yfolds, nfolds, rotation, trainsize):
    '''TODO
Determines the fold indices for the train, val, and test set given
the total number of folds, rotation, and training set size.
Use these fold indices to get the training, validation, and test sets
from all_xfolds and all_folds
    '''

    # Detrmine fold indices
    trainfolds = (np.arange(trainsize) + rotation) % nfolds
    valfold = (nfolds - 2 + rotation) % nfolds
    testfold = (valfold + 1) % nfolds

    # TODO: Construct train set by concatenating individual training
    # folds together (hint: see np.take() and np.concatenate())

    #Constructing train set
    X = np.concatenate([all_Xfolds[f] for f in trainfolds], axis=0)
    y = np.concatenate([all_yfolds[f] for f in trainfolds], axis=0)

    # TODO: Construct validation set using the valfold.
    # Hint: this is always one fold

    #Constructing validation set
    Xval = all_Xfolds[valfold]
    yval = all_yfolds[valfold]

    # TODO: Construct test set using the testfold

```



```

#Constructing test set
Xtest = all_Xfolds[testfold]
ytest = all_yfolds[testfold]

return X, y, Xval, yval, Xtest, ytest

```

```

def perform_cross_validation(self, all_Xfolds, all_yfolds, trainsize):
    ''' TODO: This is where the bulk of the work will be done
    Perform cross validation for a singular train set size and single
    hyper-parameter set, by evaluating the model's performance over
    multiple data set rotations all of the same size.

```

NOTE: This function assumes the hyper-parameters have already been set in the model

PARAMS:

all_Xfolds: list containing all of the input data folds
all_yfolds: list containing all of the output data folds
trainsize: number of folds to use for training

RETURNS: train, val, and test set results for all rotations of the data sets and the summary (i.e. the averages over all the rotations) of the results. results is a dictionary of dictionaries of r-by-n numpy arrays. Where r is the number of rotations, and n is the number of outputs from the model. summary is a dict of dictionaries of 1-by-n numpy arrays

→containing

the mean and standard deviation of the metrics in results

→across

all rotations

In our dataset, $n = 2$ (shoulder torque and elbow torque)

General form:

```
results.keys() = ['train', 'val', 'test']
```

```
results['train'].keys() = ['metric1', 'metric2', ...]
```

```
results['train']['metric1'] = numpy_array
```

```
results =
```

```
{
```

```
    'train':
```

```
    {
```

```
        'mse' : r_by_n_numpy_array,
```

```
        'rmse': r_by_n_numpy_array,
```

```
        ...
```

```

        },
        'val' : {...},
        'test' : {...}
    }

    summary =
    {
        'train':
        {
            'mse_mean' : 1_by_n_numpy_array,
            'mse_std' : 1_by_n_numpy_array,
            'rmse_mean' : 1_by_n_numpy_array,
            'rmse_std' : 1_by_n_numpy_array,
            ...
        },
        'val' : {...},
        'test' : {...}
    }

    For example, you can access the MSE results for the
    validation set like so:
        results['train'][metric]
    For example, you can access the summary (i.e. the average
    results over all the rotations) for the test set for the
    rMSE like so:
        summary['test']['rmse_mean']
'''

# Verify a valid train set size was provided
nfolds = len(all_Xfolds)
if trainsize > nfolds - 2:
    err_msg = "ERROR: KFoldHolisticCrossValidation.
↳perform_cross_validation() - "
    err_msg += "trainsize (%d) cant be more than nfolds (%d) - 2" %_
↳(trainsize, nfolds)
    raise ValueError(err_msg)

# Set up results recording for each rotation
results = {'train': None, 'val': None, 'test': None}
summary = {'train': {}, 'val': {}, 'test': {}}

model = self.model
evaluate = self.eval_func

# TODO: Rotate through different train, val, and test sets
for rotation in range(0, nfolds, self.rotation_skip):
    # Determine fold indices for train, val, and test set.

```

```

        X, y, Xval, yval, Xtest, ytest = self.get_data(all_Xfolds,
↪all_yfolds,
                                                    nfolds, rotation,
↪trainsize)

        # TODO: Train model using the training set

        #training the model with training set
        model.fit(X,y)

        # TODO: Predict with the model for train, val, and test sets

        #calling predict using train, val and test sets
        preds = model.predict(X)
        preds_val = model.predict(Xval)
        preds_test = model.predict(Xtest)

        # TODO: Evaluate the model for each set

        #calling evaluate, passing the appropriate arguments
        res_train = evaluate(model, X, y, preds)
        res_val = evaluate(model, Xval, yval, preds_val)
        res_test = evaluate(model, Xtest, ytest, preds_test)

        # Record the train, val, and test set results. These are dicts
        # of result metrics, returned by the evaluate function
        # TODO: For the first rotation, store the results from evaluating
        #         with the train, val, and tests by setting the values of
        #         the appropriate items within the results dict

        #storing results
        if results['train'] is None:
            results['train'] = res_train
            results['val'] = res_val
            results['test'] = res_test
        else:
            # Append the results for each rotation
            for metric in res_train.keys():
                results['train'][metric] = np.
↪append(results['train'][metric],
                                                    res_train[metric],
↪axis=0)

                results['val'][metric] = np.append(results['val'][metric],
                                                    res_val[metric], axis=0)
                results['test'][metric] = np.append(results['test'][metric],

```

```

res_test[metric],
axis=0)

    # Compute and record the mean and standard deviation for the given size
    for each metric
        for metric in results['train'].keys():
            for stat_set in ['train', 'val', 'test']:
                summary[stat_set][metric+'_mean'] = np.
                mean(results[stat_set][metric],
                    axis=0).reshape(1,
                -1)
                summary[stat_set][metric+'_std'] = np.
                std(results[stat_set][metric],
                    axis=0).reshape(1, -1)

    return results, summary

def grid_cross_validation(self, all_Xfolds, all_yfolds):
    ''' TODO
    (MAIN PROCEDURE) Perform cross validation for multiple sets of
    parameters and train set sizes. Calls self.perform_cross_validation().
    This is the procedure that executes cross validation for all parameter
    sets and all sizes.

    PARAMS:
        all_Xfolds: all the input data folds (list of folds, as it was
                    loaded from the files)
        all_yfolds: all the output data folds (list of folds)

    RETURNS: best parameter set for each train set size as a list of
              parameter indices. Additionally, returns self.report_by_size,
              the 3D array of validation means (overall rotations) for all
              paramsets, for each metric, for all sizes. The structure of
              the returned object is a dictionary of the following form:
              {
                  'report_by_size' : self.report_by_size,
                  'best_param_inds': self.best_param_inds
              }
    '''
    sizes = self.trainsizes
    paramsets = self.paramsets
    nparamsets = len(paramsets)
    print("nparamsets", nparamsets)

    # Set up all results
    all_results = []

```

```

# Iterate over parameter sets
for params in paramsets:
    # Set up paramset results
    param_res = []
    param_smry = None

    # Set model parameters
    print("Current paramset\n", params)
    self.model.set_params(**params)

    # Iterate over the different train set sizes
    for size in sizes:
        # TODO: Cross-validation for current model and train size
        res, smry = self.perform_cross_validation(all_Xfolds,
→all_yfolds, size)

        # Save the results
        param_res.append(res)
        # Save the mean and standard deviation statistics (summary)
        if param_smry is None:
            param_smry = smry
        else:
            # For each metric measured, append the summary results
            for metric in smry['train'].keys():
                for stat_set in ['train', 'val', 'test']:
                    stat = smry[stat_set][metric]
                    param_smry[stat_set][metric] = np.
→append(param_smry[stat_set][metric],
stat,
→axis=0)

    # Append the results and summary for the parameter set
    all_results.append({'params':params, 'results':param_res,
                        'summary':param_smry})

# Generate reports and determine best params for each size
self.results = all_results
self.report_by_size = self.get_reports()
self.best_param_inds = self.get_best_params(self.opt_metric,
self.maximize_opt_metric)

return {'report_by_size':self.report_by_size,
        'best_param_inds':self.best_param_inds}

def get_reports(self):
    ''' PROVIDED
    Get the mean validation summary of all the parameters for each size
    for all metrics. This is used to determine the best parameter set

```

```

    for each size

    RETURNS: the report_by_size as a 3D s-by-r-by-p array. Where s is
            the number of train sizes tried, r is the number of summary
            metrics evaluated+2, and p is the number of parameter sets.
    '''
    results = self.results
    sizes = np.reshape(self.trainsizes, (1, -1))

    nsizes = sizes.shape[1]
    nparams = len(results)

    # Set up the reports objects
    metrics = list(results[0]['summary']['val'].keys())
    colnames = ['params', 'size'] + metrics
    report_by_size = np.empty((nsizes, len(colnames), nparams),
    dtype=object)

    # Determine mean val for each paramset for each size for all metrics
    for p, paramset_result in enumerate(results):
        params = paramset_result['params']
        res_val = paramset_result['summary']['val']

        # Compute mean val result for each train size for each metric
        means_by_size = [np.mean(res_val[metric], axis=1) for metric in
    metrics]

        # Include the train set sizes into the report
        means_by_size = np.append(sizes, means_by_size, axis=0)
        # Include the parameter sets into the report
        param_strgs = np.reshape([str(params)]*nsizes, (1, -1))
        means_by_size = np.append(param_strgs, means_by_size, axis=0).T
        # Append the parameter set means into the report
        report_by_size[:, :, p] = means_by_size
    return report_by_size

def get_best_params(self, opt_metric, maximize_opt_metric):
    ''' PROVIDED (Do read through all the provided code)
    Determines the best parameter set for each train size, based
    on a specific metric.

    PARAMS:
        opt_metric: optimized metric. one of the metrics returned
                    from eval_func, with '_mean' appended for the
                    summary stat. This is the mean metric used to
                    determine the best parameter set for each size

        maximize_opt_metric: True if the max of opt_metric should be
    '''

```

```

        used to determine the best parameters.
        False if the min should be used.
RETURNS: list of best parameter set indicies for each size
'''
results = self.results
report_by_size = self.report_by_size

metrics = list(results[0]['summary']['val'].keys())

# Determine best params for each size, for the optimized metric
best_param_inds = None
metric_idx = metrics.index(opt_metric)

if maximize_opt_metric:
    # Add two for the additional cols for params and size
    best_param_inds = np.argmax(report_by_size[:, metric_idx+2, :],
↪axis=1)
else:
    best_param_inds = np.argmin(report_by_size[:, metric_idx+2, :],
↪axis=1)
    # Return list of best params indices for each size
    return best_param_inds

def get_best_params_strings(self):
    ''' PROVIDED
    Generates a list of strings of the best params for each size
    RETURNS: list of strings of the best params for each size
    '''
    best_param_inds = self.best_param_inds
    results = self.results
    return [str(results[p]['params']) for p in best_param_inds]

def get_report_best_params_for_size(self, size):
    ''' PROVIDED
    Get the mean validation summary for the best parameter set
    for a specific size for all metrics.
    PARAMS:
        size: index of desired train set size for the best
            paramset to come from. Size here is the index in
            the trainsizes list, NOT the actual number of folds.
    RETURNS: the best parameter report for the size as an s-by-m
            dataframe. Where each row is for a different size, and
            each column is for a different summary metric.
    '''
    best_param_inds = self.best_param_inds
    report_by_size = self.report_by_size

```

```

bp_index = best_param_inds[size]

metrics = list(self.results[0]['summary']['val'].keys())
colnames = ['params', 'size'] + metrics
report_best_params_for_size = pd.DataFrame(report_by_size[:, :, bp_index],
                                           columns=colnames)

return report_best_params_for_size

def plot_cv(self, foldsindices, results, summary, metrics, size):
    ''' PROVIDED
    Plotting function for after perform_cross_validation(),
    displaying the train and val set performances for each rotation
    of the training set.

    PARAMS:
        foldsindices: indices of the train sets tried
        results: results from perform_cross_validation()
        summary: mean and standard deviations of the results
        metrics: list of result metrics to plot. Available metrics
                are the keys in the dict returned by eval_func
        size: train set size

    RETURNS: the figure and axes handles
    '''
    nmetrics = len(metrics)

    # Initialize figure plots
    fig, axs = plt.subplots(nmetrics, 1, figsize=(12,6))
    fig.subplots_adjust(hspace=.4)
    # When 1 metric is provided, allow the axs to be iterable
    axs = np.array(axs).ravel()

    # Construct each subplot
    for metric, ax in zip(metrics, axs):
        # Compute the mean for multiple outputs
        res_train = np.mean(results['train'][metric], axis=1)
        res_val = np.mean(results['val'][metric], axis=1)
        # Plot
        ax.plot(foldsindices, res_train, label='train')
        ax.plot(foldsindices, res_val, label='val')
        ax.set(ylabel=metric)
    axs[0].legend(loc='upper right')
    axs[0].set(xlabel='Fold Index')
    axs[0].set(title='Performance for Train Set Size ' + str(size))
    return fig, axs

def plot_param_train_val(self, metrics, paramidx=0, view_test=False):

```



```

''' PROVIDED
Plotting function for after grid_cross_validation(),
displaying the mean (summary) train and val set performances
for each train set size.

PARAMS:
    metrics: list of summary metrics to plot. '_mean' or '_std'
             must be append to the end of the base metric name.
             These base metric names are the keys in the dict
             returned by eval_func
    paramidx: parameter set index
    view_test: flag to view the test set results

RETURNS: the figure and axes handles
'''

sizes = self.trainsizes
results = self.results

summary = results[paramidx]['summary']
params = results[paramidx]['params']

nmetrics = len(metrics)

# Initialize figure plots
fig, axs = plt.subplots(nmetrics, 1, figsize=(12,6))
fig.subplots_adjust(hspace=.4)
# When 1 metric is provided, allow the axs to be iterable
axs = np.array(axs).ravel()
# Construct each subplot
for metric, ax in zip(metrics, axs):
    # Compute the mean for multiple outputs
    res_train = np.mean(summary['train'][metric], axis=1)
    res_val = np.mean(summary['val'][metric], axis=1)
    # Plot
    ax.plot(sizes, res_train, label='train')
    ax.plot(sizes, res_val, label='val')
    if view_test:
        res_test = np.mean(summary['test'][metric], axis=1)
        ax.plot(sizes, res_test, label='test')
    ax.set(ylabel=metric)
axs[-1].set(xlabel='Train Set Size (# of folds)')
axs[0].set(title=str(params))
axs[0].legend(loc='upper right')
return fig, axs

def plot_allparams_val(self, metrics):
    ''' PROVIDED

```

Plotting function for after grid_cross_validation(), displaying mean (summary) validation set performances for each train size for all parameter sets for the specified metrics.

PARAMS:

metrics: list of summary metrics to plot. '_mean' or '_std' must be append to the end of the base metric name. These base metric names are the keys in the dict returned by eval_func

RETURNS: the figure and axes handles

```
'''
sizes = self.trainsizes
results = self.results

nmetrics = len(metrics)

# Initialize figure plots
fig, axs = plt.subplots(nmetrics, 1, figsize=(10,6))
fig.subplots_adjust(hspace=.4)
# When 1 metric is provided, allow the axs to be iterable
axs = np.array(axs).ravel()

# Construct each subplot
for metric, ax in zip(metrics, axs):
    for p, param_results in enumerate(results):
        summary = param_results['summary']
        params = param_results['params']
        # Compute the mean for multiple outputs
        res_val = np.mean(summary['val'][metric], axis=1)
        ax.plot(sizes, res_val, label=str(params))
    ax.set(ylabel=metric)
axs[-1].set(xlabel='Train Set Size (# of folds)')
axs[0].set(title='Validation Performance')
axs[0].legend(bbox_to_anchor=(1.02, 1), loc='upper left',
              ncol=1, borderaxespad=0., prop={'size': 8})
return fig, axs
```

```
def plot_best_params_by_size(self):
```

''' PROVIDED

Plotting function for after grid_cross_validation(), displaying mean (summary) train and validation set performances for the best parameter set for each train size for the specified metrics.

RETURNS: the figure and axes handles

```
'''
results = self.results
```

```

metric = self.opt_metric
best_param_inds = self.best_param_inds
sizes = np.array(self.trainsizes)

# Unique set of best params for the legend
unique_param_sets = np.unique(best_param_inds)
lgnd_params = [self.paramsets[p] for p in unique_param_sets]

# Initialize figure
fig, axs = plt.subplots(2, 1, figsize=(10,6))
fig.subplots_adjust(hspace=.4)
# When 1 metric is provided, allow the axs to be iterable
axs = np.array(axs).ravel()
set_names = ['train', 'val']

# Construct each subplot
for i, (ax, set_name) in enumerate(zip(axs, set_names)):
    for p in unique_param_sets:
        # Obtain indices of sizes this paramset was best for
        param_size_inds = np.where(best_param_inds == p)[0]
        param_sizes = sizes[param_size_inds]
        # Compute the mean over multiple outputs for each size
        param_summary = results[p]['summary'][set_name]
        metric_scores = np.mean(param_summary[metric][param_size_inds, :
→], axis=1)

        # Plot the param results for each size it was the best for
        ax.scatter(param_sizes, metric_scores, s=120, marker=(p+2, 1))
        #ax.grid(True)

        set_name += ' Set Performance'
        ax.set(ylabel=metric, title=set_name)

    axs[-1].set(xlabel='Train Set Size (# of folds)')
    axs[0].legend(lgnd_params, bbox_to_anchor=(1.02, 1), loc='upper left',
                  ncol=1, borderaxespad=0., prop={'size': 8})
return fig, axs

```

6 PERFORM CROSS VALIDATION FOR ELASTICNET

```

[10]: """ TODO
Generate list of parameters to use for cross validation
using generate_paramsets()
"""
param_lists = {'alpha': [.001, .005, .01, .05, .1],
               'l1_ratio': [.05, .1], 'max_iter': [1e4]}

```

```
#calling generate_paramsets()
allparamsets = generate_paramsets(param_lists)
allparamsets
```

```
[10]: [{'alpha': 0.001, 'l1_ratio': 0.05, 'max_iter': 10000.0},
      {'alpha': 0.001, 'l1_ratio': 0.1, 'max_iter': 10000.0},
      {'alpha': 0.005, 'l1_ratio': 0.05, 'max_iter': 10000.0},
      {'alpha': 0.005, 'l1_ratio': 0.1, 'max_iter': 10000.0},
      {'alpha': 0.01, 'l1_ratio': 0.05, 'max_iter': 10000.0},
      {'alpha': 0.01, 'l1_ratio': 0.1, 'max_iter': 10000.0},
      {'alpha': 0.05, 'l1_ratio': 0.05, 'max_iter': 10000.0},
      {'alpha': 0.05, 'l1_ratio': 0.1, 'max_iter': 10000.0},
      {'alpha': 0.1, 'l1_ratio': 0.05, 'max_iter': 10000.0},
      {'alpha': 0.1, 'l1_ratio': 0.1, 'max_iter': 10000.0}]
```

```
[11]: """ TODO
Initialize the cross validation object. Use ElasticNet for the
ase model, use every even value between 2 and 18, inclusive, for
the train set sizes, use score_eval as the eval_func, use rmse
as the metric to optimize, and 4 for the skip. We want ot minimize
rmse thus set maximize_opt_metric=False
"""

#setting model to elasticnet
model = ElasticNet()
#setting trainsize to be 1-6 inclusive
trainsizes = [i for i in range(1,7)]
#setting opt_metric to 'rmse'
opt_metric = 'rmse'
#setting maximize_opt_metric to false
maximize_opt_metric = False
#setting skip to 4
skip = 4
#initializing a cross validation object
crossval = KFoldHolisticCrossValidation(model, allparamsets, score_eval,
    ↪opt_metric,
    maximize_opt_metric, trainsizes, skip)
```

```
[12]: """ TODO
Execute the grid_cross_validation() procedure for all parameters
and train set sizes
"""

# TODO: make sure this is set appropriately. True if you want to
#       just always to run cross validation, false if you want
#       to re-load a previous run

#set force to false to it can read previous run for if statement below
force = False
```

```

fullcvfname = "hw6_crossval.pkl"

crossval_report = None
if force or (not os.path.exists(fullcvfname)):
    # TODO: Use grid_cross_validation() to run the full cross
    #         validation procedure
    # Note: when testing, run this using small lists of parameters
    #       (e.g. of length 2 or 4) and/or small trainsize lists
    #       (e.g. [1, 2, 3, 4, 5])
    # Note: for the final submission, make sure to use the complete
    #       parameter set list and trainsize list provided/specified
    #       This will take some time.

    #passing MI_folds for allXfolds and torque_folds for allyfolds
    crossval_report = crossval.grid_cross_validation(MI_folds, torque_folds)
    joblib.dump(crossval, fullcvfname)
else:
    # TODO: Re-load saved crossval object instead of re-running the
    #       cross validation procedure. Use joblib.load()

    #calling joblib.load()
    crossval = joblib.load(fullcvfname)
    crossval_report = {'report_by_size' : crossval.report_by_size,
                       'best_param_inds': crossval.best_param_inds}

crossval_report.keys()

```

```
[12]: dict_keys(['report_by_size', 'best_param_inds'])
```

7 RESULTS

```

[13]: """ TODO
      Obtain all the results for all parameters, for all sizes, for all
      rotations. This is the results attribute of the crossval object
      """

      #calling crossval.results
      all_results = crossval.results
      len(all_results)

```

```
[13]: 10
```

```

[14]: """ PROVIDED
      Display the keys of the results object
      """

      all_results[0].keys()

```

```
[14]: dict_keys(['params', 'results', 'summary'])
```

```
[15]: """ TODO
Obtain and display the indices of the best parameters for each
size using either the best_params_inds attribute of the crossval
object or 'best_param_inds' item from the crossval_report dict
"""

#calling crossval_report item best_param_inds
best_param_inds = crossval_report['best_param_inds']
best_param_inds
```

```
[15]: array([5, 5, 4, 4, 4, 4])
```

```
[16]: """ TODO
Display the list of the best parameter sets for each size. Use
crossval.get_best_params_strings()
"""

# TODO

#calling crossval.get_best_params_strings()
crossval.get_best_params_strings()
```

```
[16]: [{"alpha": 0.01, 'l1_ratio': 0.1, 'max_iter': 10000.0},
{"alpha": 0.01, 'l1_ratio': 0.1, 'max_iter': 10000.0},
{"alpha": 0.01, 'l1_ratio': 0.05, 'max_iter': 10000.0},
{"alpha": 0.01, 'l1_ratio': 0.05, 'max_iter': 10000.0},
{"alpha": 0.01, 'l1_ratio': 0.05, 'max_iter': 10000.0},
{"alpha": 0.01, 'l1_ratio': 0.05, 'max_iter': 10000.0}]
```

```
[17]: """ TODO
Obtain and display the shape of the report of all the parameters'
mean results over all sizes and rotations. This is the report_by_size
attribute of the crossval object. It is also stored within the
'report_by_size' item of the crossval_report dict
"""

#calling crossval_report item report_by_size
report = crossval_report['report_by_size']
report.shape
```

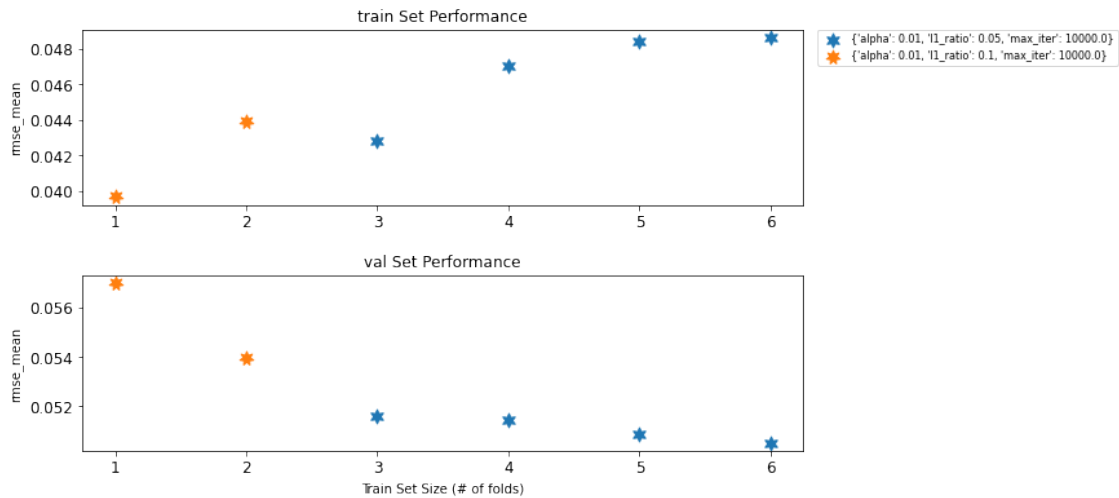
```
[17]: (6, 10, 10)
```

```
[18]: """ TODO
Plot the mean (summary) train and validation set performances for
the best parameter set for each train size for the optimized
metrics. Use plot_best_params_by_size()
"""
```

```
# TODO

#calling plot_best_params_by_size function of crossval
crossval.plot_best_params_by_size()
```

```
[18]: (<Figure size 720x432 with 2 Axes>,
       array([<AxesSubplot:title={'center':'train Set Performance'},
              ylabel='rmse_mean'>,
              <AxesSubplot:title={'center':'val Set Performance'}, xlabel='Train Set
              Size (# of folds)', ylabel='rmse_mean'>],
              dtype=object))
```



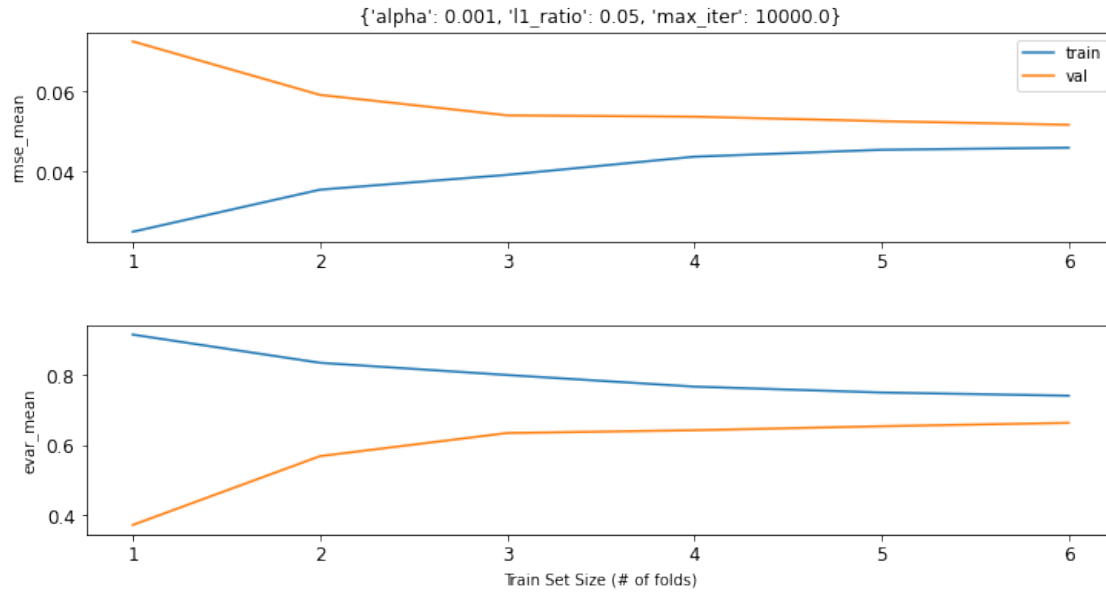
```
[25]: """ TODO
Plot the average results (summary) over train set size for all
parameter sets for the metrics 'rmse_mean' and 'evan_mean'
for the train and val sets. Use plot_param_train_val().
view_test=False
"""

metrics = ['rmse_mean', 'evan_mean']

# TODO

#calling plot_param_train_val function of crossval
crossval.plot_param_train_val(metrics, view_test=False)
```


```
[25]: (<Figure size 864x432 with 2 Axes>,
       array([<AxesSubplot:title={'center':"'alpha': 0.001, 'l1_ratio': 0.05,
              'max_iter': 10000.0}"}, ylabel='rmse_mean'>,
              <AxesSubplot:xlabel='Train Set Size (# of folds)', ylabel='evan_mean'>],
              dtype=object))
```



```
[20]: """ TODO
Plot the validation results for all parameters over all train
sizes, for the specified metrics. Use plot_allparams_val()
"""
# TODO

#calling plot_allparams_val function of crossval
crossval.plot_allparams_val(metrics)
```

```
[20]: (<Figure size 720x432 with 2 Axes>,
array([<AxesSubplot:title={'center':'Validation Performance'},
ylabel='rmse_mean'>,
<AxesSubplot:xlabel='Train Set Size (# of folds)', ylabel='evar_mean'>],
dtype=object))
```





```
[27]: """ TODO
For the best parameter set for the train set size at index 5,
plot the TRAIN, VAL, and TEST set performances using
plot_param_train_val() for just the optimized metric
"""

size_idx = 5
# TODO

#calling plot_param_train_val function of crossval
crossval.plot_param_train_val([crossval.opt_metric], best_param_inds[size_idx],
↪view_test=True)
```

```
[27]: (<Figure size 864x432 with 1 Axes>,
array([<AxesSubplot:title={'center': '{"alpha': 0.01, 'l1_ratio': 0.05,
'max_iter': 10000.0}"}, xlabel='Train Set Size (# of folds)',
ylabel='rmse_mean'>],
dtype=object))
```



```
[103]: """ TODO
Use get_report_best_params_for_size() to display the report of
the average val statistics for the best parameter set, for the
train set size at index 5 (i.e. size_idx)
"""

#calling get_report_best_params_for_size function of crossval
report_best_params = crossval.get_report_best_params_for_size(size_idx)
report_best_params
```

```
[103]:
```

	params	size	\
0	{'alpha': 0.01, 'l1_ratio': 0.05, 'max_iter': ...}	1.0	
1	{'alpha': 0.01, 'l1_ratio': 0.05, 'max_iter': ...}	2.0	
2	{'alpha': 0.01, 'l1_ratio': 0.05, 'max_iter': ...}	3.0	
3	{'alpha': 0.01, 'l1_ratio': 0.05, 'max_iter': ...}	4.0	
4	{'alpha': 0.01, 'l1_ratio': 0.05, 'max_iter': ...}	5.0	
5	{'alpha': 0.01, 'l1_ratio': 0.05, 'max_iter': ...}	6.0	

	mse_mean	mse_std	rmse_mean	\
0	0.004339758916158861	0.001695959737335727	0.05837819429918903	
1	0.0036484504643108846	0.0014904783428125472	0.05397125419285854	
2	0.003302651721721235	0.0013097047039709432	0.05158260020754825	
3	0.003288324571939147	0.0012794376602931277	0.05143339933709775	
4	0.003189258626079231	0.0012111260230939665	0.05084344419067663	
5	0.0031420698497129037	0.001230751306487161	0.05051109258741546	

	rmse_std	evvar_mean	evvar_std	\
0				
1				
2				
3				
4				
5				

0	0.01175946265587988	0.6002958533896176	0.05819936529463857
1	0.011431333558763551	0.6382307158738958	0.04986795249206236
2	0.01075893609232094	0.6622754558955954	0.04511540687545403
3	0.010594508017117338	0.665650750956378	0.04291078087335755
4	0.010209584282787808	0.6689877221669576	0.040154616347059365
5	0.01025505282563732	0.6714990587891585	0.04093577025092407

	score_mean	score_std
0	0.5964794766688885	0.061509748501615405
1	0.6362657732004847	0.05014628130905605
2	0.660758581632639	0.04449807384256146
3	0.6641574899843384	0.04261380945242854
4	0.6678621132901451	0.04007599471243523
5	0.6700803391597506	0.04017735574057148

[]: