

homework7-skel

November 11, 2020

NAME: Nigel Mansell

SECTION: 995

CS 5970: Machine Learning Practices

1 Homework 7: Model Comparisons

1.1 Assignment Overview

Generally, it's helpful to first read through the entire notebook before writing any code to obtain a sense of the overall program structure before you start coding.

Follow the TODOs and read through and understand any provided code.

1.1.1 Task

For this assignment, you'll be comparing different models after performing holistic cross validation to find the best parameter sets for various sizes of the training data.

For this assignment, we will try to predict shoulder and elbow torque simultaneously, from the neural activation.

1.1.2 Data set

The BMI data will be utilized. Recall:

* *MI* files contain data with the number of spikes for 48 neurons, at multiple time points, for a single fold. There are 20 folds (20 files), where each fold consists of over 1000 time points (the rows). At each time point, we record the number of spikes for each neuron for 20 bins. Therefore, each time point has $48 * 20 = 960$ columns.

* *theta* files record the angular position of the shoulder (in column 0) and the elbow (in column 1) for each time point. (rad) * *dtheta* files record the angular velocity of the shoulder (in column 0) and the elbow (in column 1) for each time point. (rad/s) * *torque* files record the torque of the shoulder (in column 0) and the elbow (in column 1) for each time point. (N-m) * *time* files record the actual time stamp of each time point.

1.1.3 Objectives

- Understanding regularization using **holistic cross validation**
- Training set size sensitivity analysis
- Model selection

1.1.4 Notes

- Do not save work within the ml_practices folder

1.1.5 General References

- [Guide to Jupyter](#)
- [Python Built-in Functions](#)
- [Python Data Structures](#)
- [Numpy Reference](#)
- [Numpy Cheat Sheet](#)
- [Summary of matplotlib](#)
- [DataCamp: Matplotlib](#)
- [Pandas DataFrames](#)
- [Sci-kit Learn Linear Models](#)
- [Sci-kit Learn Ensemble Models](#)
- [Sci-kit Learn Metrics](#)
- [Sci-kit Learn Model Selection](#)
- [SciPy Paired t-test for Dependent Samples](#)
- [Student's t-test](#)
- [Understanding Paired t-tests](#)

1.1.6 Hand-In Procedure

- Execute all cells so they are showing correct results
- Notebook:
 - Submit this file (.ipynb) to the Canvas HW7 dropbox
- PDF:
 - File/Print/Print to file -> Produces a copy of the notebook in PDF format
 - Submit the PDF file to the Gradescope HW7 dropbox

```
[1]: import pandas as pd
import numpy as np
import scipy.stats as stats
import os, re, fnmatch
import pathlib, itertools, time
import matplotlib.pyplot as plt
import joblib

from matplotlib import cm
from mpl_toolkits.mplot3d import Axes3D
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.metrics import explained_variance_score
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet

FIGW = 10
FIGH = 6
FONTSIZE = 12
```

```
HOME_DIR = pathlib.Path.home()

plt.rcParams['figure.figsize'] = (FIGW, FIGH)
plt.rcParams['font.size'] = FONTSIZE

plt.rcParams['xtick.labelsize'] = FONTSIZE
plt.rcParams['ytick.labelsize'] = FONTSIZE

%matplotlib inline
```

```
[2]: """
    Display current working directory of this notebook. If you are using
    relative paths for your data, then it needs to be relative to the CWD.
    """
    pathlib.Path.cwd()
```

```
[2]: PosixPath('/home/nigel/Desktop/mlp/hw7')
```

2 LOAD DATA

```
[3]: def read_bmi_file_set(directory, filebase):
    """
    Read a set of CSV files and append them together
    :param directory: The directory in which to scan for the CSV files
    :param filebase: File specification potentially including wildcards
    :returns: A list of Numpy arrays (one for each fold)
    """

    # The set of files in the directory
    files = fnmatch.filter(os.listdir(directory), filebase)
    files.sort()

    # Create list of Pandas objects;
    # Each from a file in the directory matching the filebase
    lst = [pd.read_csv(directory + "/" + file, delim_whitespace=True).values
            for file in files]

    # Concatenate the Pandas objects together. ignore_index is
    # critical here so that the duplicate row indices are addressed
    return lst
```

```
[4]: """ PROVIDED
    Load the BMI data from all the folds, using read_bmi_file_set()
    """
    # TODO: might need to change directory
```

```

dir_name = '/home/nigel/Desktop/mlp/mlp_2020/datasets/bmi/DAT6_08'

MI_folds = read_bmi_file_set(dir_name, 'MI_fold*')
theta_folds = read_bmi_file_set(dir_name, 'theta_fold*')
dtheta_folds = read_bmi_file_set(dir_name, 'dtheta_fold*')
torque_folds = read_bmi_file_set(dir_name, 'torque_fold*')
time_folds = read_bmi_file_set(dir_name, 'time_fold*')

alldata_folds = zip(MI_folds, theta_folds, dtheta_folds,
                    torque_folds, time_folds)

nfolders = len(MI_folds)
nfolders

```

[4]: 20

```

[5]: """ PROVIDED
      Print out the shape of all the data for each fold
      """
      for i, (MI, theta, dtheta, torque, time) in enumerate(alldata_folds):
          print("FOLD %2d " % i, MI.shape, theta.shape,
                dtheta.shape, torque.shape, time.shape)

```

```

FOLD 0 (1193, 960) (1193, 2) (1193, 2) (1193, 2) (1193, 1)
FOLD 1 (1104, 960) (1104, 2) (1104, 2) (1104, 2) (1104, 1)
FOLD 2 (1531, 960) (1531, 2) (1531, 2) (1531, 2) (1531, 1)
FOLD 3 (1265, 960) (1265, 2) (1265, 2) (1265, 2) (1265, 1)
FOLD 4 (1498, 960) (1498, 2) (1498, 2) (1498, 2) (1498, 1)
FOLD 5 (1252, 960) (1252, 2) (1252, 2) (1252, 2) (1252, 1)
FOLD 6 (1375, 960) (1375, 2) (1375, 2) (1375, 2) (1375, 1)
FOLD 7 (1130, 960) (1130, 2) (1130, 2) (1130, 2) (1130, 1)
FOLD 8 (1247, 960) (1247, 2) (1247, 2) (1247, 2) (1247, 1)
FOLD 9 (1257, 960) (1257, 2) (1257, 2) (1257, 2) (1257, 1)
FOLD 10 (1265, 960) (1265, 2) (1265, 2) (1265, 2) (1265, 1)
FOLD 11 (1146, 960) (1146, 2) (1146, 2) (1146, 2) (1146, 1)
FOLD 12 (1225, 960) (1225, 2) (1225, 2) (1225, 2) (1225, 1)
FOLD 13 (1238, 960) (1238, 2) (1238, 2) (1238, 2) (1238, 1)
FOLD 14 (1570, 960) (1570, 2) (1570, 2) (1570, 2) (1570, 1)
FOLD 15 (1359, 960) (1359, 2) (1359, 2) (1359, 2) (1359, 1)
FOLD 16 (1579, 960) (1579, 2) (1579, 2) (1579, 2) (1579, 1)
FOLD 17 (1364, 960) (1364, 2) (1364, 2) (1364, 2) (1364, 1)
FOLD 18 (1389, 960) (1389, 2) (1389, 2) (1389, 2) (1389, 1)
FOLD 19 (1289, 960) (1289, 2) (1289, 2) (1289, 2) (1289, 1)

```

3 PARAMETER SET LIST

```
[6]: def generate_paramsets(param_lists):  
    '''  
    Construct the Cartesian product of the parameters  
    PARAMS:  
        params_lists: dict of lists of values to try for each parameter.  
            keys of the dict are the names of the parameters  
            values are lists of values to try for the  
                corresponding parameter  
    RETURNS: a list of dicts that make up the Cartesian product of the  
        parameters  
    '''  
    keys, values = zip(*param_lists.items())  
    # Determines cartesian product of parameter values  
    combos = itertools.product(*values)  
    # Constructs list of dictionaries  
    combos_dicts = [dict(zip(keys, vals)) for vals in combos]  
    return list(combos_dicts)
```

4 PERFORMANCE EVALUATION

```
[7]: def mse_rmse(trues, preds):  
    '''  
    Compute MSE and rMSE for each column separately.  
    '''  
    mse = np.sum(np.square(trues - preds), axis=0) / trues.shape[0]  
    rmse = np.sqrt(mse)  
    return mse, rmse  
  
def score_eval(model, X, y, preds):  
    '''  
    Compute the model predictions and corresponding scores, for an  
    already trained model.  
    PARAMS:  
        model: model to predict with  
        X: input feature data  
        y: true output for X  
        preds: predicted output for X  
    RETURNS: results as a dictionary of numpy arrays  
        mse: mean squared error for each column  
        rmse: rMSE for each column  
        evar: explained variance, best is 1.0  
        score: score computed by the models score() method  
    '''  
    score = model.score(X, y)
```

```

mse, rmse = mse_rmse(y, preds)
evar = explained_variance_score(y, preds)

# Dictionary of numpy arrays. The numpy arrays must
# be row vectors, where each element is the result
# for a different output, when using multiple regression.
# The keys of the dictionary are the name of the performance
# metric, and the values are the numpy row vectors
results = {'mse': np.reshape(mse, (1, -1)),
           'rmse': np.reshape(rmse, (1, -1)),
           'evar': np.reshape(evar, (1, -1)),
           'score': np.reshape(score, (1, -1)),
          }
return results

```

5 CROSS VALIDATION

```

[8]: """ PROVIDED:
      This is the same KFoldHolisticCrossValidation class from HW6.
      """
class KFoldHolisticCrossValidation():
    def __init__(self, model, paramsets, eval_func, opt_metric,
                  maximize_opt_metric=False, trainsizes=[1], rotation_skip=1):
        """
        Object for managing and performing cross validation for a given
        model for a list of parameter sets and train set sizes. Note,
        train set size is in terms of number of folds (not samples)

        General Procedure:
        + iter over hyper-parameter sets
          1. set hyper-parameters of the model
          2. iter over train set sizes
             a. iter over splits/rotations
                i. train the model
                ii. evaluate the model on train, val, and test sets
                iii. record the results
             b. record the results by size
          3. record the results by hyper-parameter set

        PARAMS:
            model: base ML model

            paramsets: list of dicts of parameter sets to give to the model

            eval_func: handle to function used to evaluate/score the model

```

```

        The eval_func definition must have the following
        arguments: model, X, ytrue, ypreds; and return a dict
        of numpy arrays with shape 1-by-n, where n is the
        number of outputs if using multiple regression.
        template function header:
            def eval_func(model, X, y, preds)
        template output:
            {'metrics1':1_by_n_array, ...}

    opt_metric: the optimized metric. one of the metric key names
                returned from eval_func to use to pick the best
                parameter sets

    maximize_opt_metric: True if opt_metric is maximized;
                        False if minimized

    trainsizes: list of training set sizes (in number of folds) to try

    rotation_skip: build model and evaluate every ith rotation (1=all
                    possible rotations; 2=every other rotation, etc.)
'''
self.model = model
self.paramsets = paramsets
self.trainsizes = trainsizes
self.eval_func = eval_func
self.opt_metric = opt_metric + '_mean'
self.maximize_opt_metric = maximize_opt_metric
self.rotation_skip = rotation_skip

# Results attributes
# Full recording of all results for all paramsets, sizes, rotations,
# and metrics. This is a list of dictionaries for each paramset
self.results = None
# Validation summary report of all means and standard deviations for
# all metrics, for all paramsets, and sizes. This is a 3D s-by-r-by-p
# numpy array. Where s is the number of sizes, r the number of summary
# metrics +2, and p is the number of paramsets
self.report_by_size = None
# List of the indices of the best paramset for each size
self.best_param_inds = None

def get_data(self, all_Xfolds, all_yfolds, nfolds, rotation, trainsize):
    '''
    Determines the fold indices for the train, val, and test set given
    the total number of folds, rotation, and training set size.
    Use these fold indices to get the training, validation, and test sets

```

```

from all_xfolds and all_folds
'''
# Determine fold indices
trainfolds = (np.arange(trainsize) + rotation) % nfolds
valfold = (nfolds - 2 + rotation) % nfolds
testfold = (valfold + 1) % nfolds

# Construct train set by concatenating the individual
#     training folds together (hint: see np.take() and
#     np.concatenate())
X = np.concatenate(np.take(all_Xfolds, trainfolds))
y = np.concatenate(np.take(all_yfolds, trainfolds))

# Construct validation set. Hint: this is always one fold
Xval = all_Xfolds[valfold]
yval = all_yfolds[valfold]

# Construct test set
Xtest = all_Xfolds[testfold]
ytest = all_yfolds[testfold]

return X, y, Xval, yval, Xtest, ytest

def perform_cross_validation(self, all_Xfolds, all_yfolds,
                             trainsize):
    '''
    Perform cross validation for a singular train set size and single
    hyper-parameter set, by evaluating the model's performance over
    multiple data set rotations all of the same size.

    NOTE: This function assumes the hyper-parameters have already been
           set in the model

    PARAMS:
        all_Xfolds: list containing all of the input data folds
        all_yfolds: list containing all of the output data folds
        trainsize: number of folds to use for training

    RETURNS: train, val, and test set results for all rotations of the
              data sets and the summary (i.e. the averages over all the
              rotations) of the results.
              results is a dictionary of dictionaries of r-by-n numpy
              arrays. Where r is the number of rotations, and n is the
              number of outputs from the model.
              summary is a dictionary of dictionaries of 1-by-n numpy
              arrays.

```


General form:

```
results.keys() = ['train', 'val', 'test']

results['train'].keys() = ['metric1', 'metric2', ...]

results['train']['metric1'] = numpy_array

results =
{
    'train':
        {
            'mse'      : r_by_n_numpy_array,
            'rmse': r_by_n_numpy_array,
            ...
        },
    'val'  : {...},
    'test' : {...}
}

summary =
{
    'train':
        {
            'mse_mean'      : 1_by_n_numpy_array,
            'mse_std'       : 1_by_n_numpy_array,
            'rmse_mean': 1_by_n_numpy_array,
            'rmse_std' : 1_by_n_numpy_array,
            ...
        },
    'val'  : {...},
    'test' : {...}
}
```

For example, you can access the MSE results for the validation set like so:

```
results['val']['mse']
```

For example, you can access the summary (i.e. the average results over all the rotations) for the test set for the rMSE like so:

```
summary['test']['rmse_mean']
```

```
...
```

```
# Verify a valid train set size was provided
```

```
nfolds = len(all_Xfolds)
```

```
if trainsize < 1 or trainsize > nfolds - 2:
```

```

        err_msg = "ERROR: KFoldHolisticCrossValidation.
→perform_cross_validation() - "
        err_msg += "trainsize (%d) must be between 1 and nfolds (%d) - 2" % (
→(trainsize, nfolds)
            raise ValueError(err_msg)

        # Verify rotation skip
        if self.rotation_skip < 1:
            err_msg = "ERROR: KFoldHolisticCrossValidation.__init__() - "
            err_msg += "rotation_skip (%d) can't be less than 1" % self.
→rotation_skip
            raise ValueError(err_msg)

        # Set up results recording for each rotation
        results = {'train': None, 'val': None, 'test': None}
        summary = {'train': {}, 'val': {}, 'test': {}}

        model = self.model
        evaluate = self.eval_func

        print(range(0, nfolds, self.rotation_skip))

        # Rotate through different train, val, and test sets
        for rotation in range(0, nfolds, self.rotation_skip):
            # Determine fold indices for train, val, and test set.
            X, y, Xval, yval, Xtest, ytest = self.get_data(all_Xfolds,
→all_yfolds,
                                                    nfolds, rotation,
→trainsize)

            # Train model using the training set
            model.fit(X, y) # make sure warm_start is False

            # Predict with the model for train, val, and test sets
            preds = model.predict(X)
            preds_val = model.predict(Xval)
            preds_test = model.predict(Xtest)

            # Evaluate the model for each set
            res_train = evaluate(model, X, y, preds)
            res_val = evaluate(model, Xval, yval, preds_val)
            res_test = evaluate(model, Xtest, ytest, preds_test)

            # Record the train, val, and test set results. These are dicts
            # of result metrics, returned by the evaluate function
            # For the first rotation, store the results from evaluating
            # with the train, val, and tests by setting the values of

```

```

        # the appropriate items within the results dict
    if results['train'] is None:
        results['train'] = res_train
        results['val'] = res_val
        results['test'] = res_test
    else:
        # Append the results for each rotation
        for metric in res_train.keys():
            results['train'][metric] = np.
→append(results['train'][metric],
                                                res_train[metric],
→axis=0)

            results['val'][metric] = np.append(results['val'][metric],
                                                res_val[metric], axis=0)
            results['test'][metric] = np.append(results['test'][metric],
                                                res_test[metric],
→axis=0)

        # Compute/record mean and standard deviation for the size for each
→metric
        for metric in results['train'].keys():
            for stat_set in ['train', 'val', 'test']:
                summary[stat_set][metric+'_mean'] = np.
→mean(results[stat_set][metric],
                                                axis=0).reshape(1,
→-1)

                summary[stat_set][metric+'_std'] = np.
→std(results[stat_set][metric],
                                                axis=0).reshape(1, -1)

    return results, summary

def grid_cross_validation(self, all_Xfolds, all_yfolds):
    """
    (MAIN PROCEDURE) Perform cross validation for multiple sets of
    parameters and train set sizes. Calls self.perform_cross_validation().
    This is the procedure that executes cross validation for all parameter
    sets and all sizes.

    General Procedure:
    + iter over hyper-parameter sets
      1. set hyper-parameters of the model
      2. iter over train set sizes
         a. iter over splits/rotations
            i. train the model
            ii. evaluate the model on train, val, and test sets
    """

```

```

        iii. record the results
    b. record the results by size
    3. record the results by hyper-parameter set

PARAMS:
    all_Xfolds: all the input data folds (list of folds, as it was
        loaded from the files)
    all_yfolds: all the output data folds (list of folds)

RETURNS: best parameter set for each train set size as a list of
        parameter indices. Additionally, returns self.report_by_size,
        the 3D array of validation means (overall rotations) for all
        paramsets, for each metric, for all sizes. The structure of
        the returned object is a dictionary of the following form:
    {
        'report_by_size' : self.report_by_size,
        'best_param_inds': self.best_param_inds
    }
    '''

    sizes = self.trainsizes
    paramsets = self.paramsets
    nparamsets = len(paramsets)
    print("nparamsets", nparamsets)

    # Set up all results
    all_results = []

    # Iterate over parameter sets
    for params in paramsets:
        # Set up paramset results
        param_res = []
        param_smry = None

        # Set model parameters
        print("Current paramset\n", params)
        self.model.set_params(**params)
        print('Sizes:', sizes)
        # Iterate over the different train set sizes
        for size in sizes:
            print('Current size:', size)
            # Cross-validation for current model and train size
            res, smry = self.perform_cross_validation(all_Xfolds,
                                                        all_yfolds,
                                                        size)

            # Save the results
            param_res.append(res)

```

```

        # Save the mean and standard deviation statistics (summary)
        if param_smry is None: param_smry = smry
        else:
            # For each metric measured, append the summary results
            for metric in smry['train'].keys():
                for stat_set in ['train', 'val', 'test']:
                    stat = smry[stat_set][metric]
                    param_smry[stat_set][metric] = np.
→append(param_smry[stat_set][metric],
                                                    stat,
→axis=0)

        # Append the results and summary for the parameter set
        all_results.append({'params':params, 'results':param_res,
                            'summary':param_smry})

        # Generate reports and determine best params for each size
        self.results = all_results
        self.report_by_size = self.get_reports()
        self.best_param_inds = self.get_best_params(self.opt_metric,
                                                    self.maximize_opt_metric)

        return {'report_by_size':self.report_by_size,
                'best_param_inds':self.best_param_inds}

def get_reports(self):
    """
    Get the mean validation summary of all the parameters for each size
    for all metrics. This is used to determine the best parameter set
    for each size

    RETURNS: the report_by_size as a 3D s-by-r-by-p array. Where s is
    the number of train sizes tried, r is the number of summary
    metrics evaluated+2, and p is the number of parameter sets.
    """
    results = self.results
    sizes = np.reshape(self.trainsizes, (1, -1))

    nsizes = sizes.shape[1]
    nparams = len(results)

    # Set up the reports objects
    metrics = list(results[0]['summary']['val'].keys())
    colnames = ['params', 'size'] + metrics
    report_by_size = np.empty((nsizes, len(colnames), nparams), dtype=object)

    # Determine mean val for each paramset for each size for all metrics
    for p, paramset_result in enumerate(results):

```

```

params = paramset_result['params']
res_val = paramset_result['summary']['val']

# Compute mean val result for each train size for each metric
means_by_size = [np.mean(res_val[metric], axis=1)
                  for metric in metrics]
# Include the train set sizes into the report
means_by_size = np.append(sizes, means_by_size, axis=0)
# Include the parameter sets into the report
param_strgs = np.reshape([str(params)]*nsizes, (1, -1))
means_by_size = np.append(param_strgs, means_by_size, axis=0).T
# Append the parameter set means into the report
report_by_size[:, :, p] = means_by_size
return report_by_size

def get_best_params(self, opt_metric, maximize_opt_metric):
    '''
    Determines the best parameter set for each train size,
    based on a specific metric.

    PARAMS:
        opt_metric: optimized metric. one of the metrics returned
                    from eval_func, with '_mean' appended for the
                    summary stat. This is the mean metric used to
                    determine the best parameter set for each size

        maximize_opt_metric: True if the max of opt_metric should be
                             used to determine the best parameters.
                             False if the min should be used.

    RETURNS: list of best parameter set indicies for each size
    '''
    results = self.results
    report_by_size = self.report_by_size

    metrics = list(results[0]['summary']['val'].keys())

    # Determine best params for each size, for the optimized metric
    best_param_inds = None
    metric_idx = metrics.index(opt_metric)

    # Report info for all paramsets for the optimized metric
    report_opt_metric = report_by_size[:, metric_idx+2, :]

    if maximize_opt_metric:
        # Add two for the additional cols for params and size
        best_param_inds = np.argmax(report_opt_metric, axis=1)
    else:

```

```

        best_param_inds = np.argmin(report_opt_metric, axis=1)
        # Return list of best params indices for each size
        return best_param_inds

def get_best_params_strings(self):
    """
    Generates a list of strings of the best params for each size
    RETURNS: list of strings of the best params for each size
    """
    best_param_inds = self.best_param_inds
    results = self.results
    return [str(results[p]['params']) for p in best_param_inds]

def get_report_best_params_for_size(self, size):
    """
    Get the mean validation summary for the best parameter set
    for a specific size for all metrics.
    PARAMS:
        size: index of desired train set size for the best
              paramset to come from. Size here is the index in
              the trainsizes list, NOT the actual number of folds.
    RETURNS: the best parameter report for the size as an s-by-m
              dataframe. Where each row is for a different size, and
              each column is for a different summary metric.
    """
    best_param_inds = self.best_param_inds
    report_by_size = self.report_by_size

    # Obtain the index of the best parameter set for the size
    bp_index = best_param_inds[size]

    # Obtain the list of metrics
    metrics = list(self.results[0]['summary']['val'].keys())
    colnames = ['params', 'size'] + metrics

    # Create DataFrame with all summary stats for the parameter set
    report_best_params_for_size = pd.DataFrame(report_by_size[:, :, bp_index],
                                                columns=colnames)

    return report_best_params_for_size

def plot_cv(self, foldsindices, results, summary, metrics, size):
    """
    Plotting function for after perform_cross_validation(),
    displaying the train and val set performances for each rotation
    of the training set.

    PARAMS:

```

```

        foldsindices: indices of the train sets tried
        results: results from perform_cross_validation()
        summary: mean and standard deviations of the results
        metrics: list of result metrics to plot. Available metrics
                  are the keys in the dict returned by eval_func
        size: train set size

    RETURNS: the figure and axes handles
    '''
    nmetrics = len(metrics)

    # Initialize figure plots
    fig, axs = plt.subplots(nmetrics, 1, figsize=(12,6))
    fig.subplots_adjust(hspace=.35)
    # When 1 metric is provided, allow the axs to be iterable
    axs = np.array(axs).ravel()

    # Construct each subplot
    for metric, ax in zip(metrics, axs):
        # Compute the mean for multiple outputs
        res_train = np.mean(results['train'][metric], axis=1)
        res_val = np.mean(results['val'][metric], axis=1)
        #res_test = np.mean(results['test'][metric], axis=1)
        # Plot
        ax.plot(foldsindices, res_train, label='train')
        ax.plot(foldsindices, res_val, label='val')
        #ax.plot(foldsindices, res_test, label='test')
        ax.set ylabel=metric)
    axs[0].legend(loc='upper right')
    axs[0].set xlabel='Fold Index')
    axs[0].set title='Performance for Train Set Size ' + str(size))
    return fig, axs

def plot_param_train_val(self, metrics, paramidx=0, view_test=False):
    '''
    Plotting function for after grid_cross_validation(),
    displaying the mean (summary) train and val set performances
    for each train set size.

    PARAMS:
        metrics: list of summary metrics to plot. '_mean' or '_std'
                  must be append to the end of the base metric name.
                  These base metric names are the keys in the dict
                  returned by eval_func
        paramidx: parameter set index
        view_test: flag to view the test set results
    '''

```



```

RETURNS: the figure and axes handles
'''
    sizes = self.trainsizes
    results = self.results

    summary = results[paramidx]['summary']
    params = results[paramidx]['params']

    nmetrics = len(metrics)

    # Initialize figure plots
    fig, axs = plt.subplots(nmetrics, 1, figsize=(12,6))
    fig.subplots_adjust(hspace=.35)
    # When 1 metric is provided, allow the axs to be iterable
    axs = np.array(axs).ravel()

    # Construct each subplot
    for metric, ax in zip(metrics, axs):
        # Compute the mean for multiple outputs
        res_train = np.mean(summary['train'][metric], axis=1)
        res_val = np.mean(summary['val'][metric], axis=1)
        # Plot
        ax.plot(sizes, res_train, label='train')
        ax.plot(sizes, res_val, label='val')
        if view_test:
            res_test = np.mean(summary['test'][metric], axis=1)
            ax.plot(sizes, res_test, label='test')
        ax.set(ylabel=metric)
    axs[-1].set(xlabel='Train Set Size (# of folds)')
    axs[0].set(title=str(params))
    axs[0].legend(loc='upper right')
    return fig, axs

def plot_allparams_val(self, metrics):
    '''
    Plotting function for after grid_cross_validation(), displaying
    mean (summary) validation set performances for each train size
    for all parameter sets for the specified metrics.

    PARAMS:
        metrics: list of summary metrics to plot. '_mean' or '_std'
        must be append to the end of the base metric name.
        These base metric names are the keys in the dict
        returned by eval_func

    RETURNS: the figure and axes handles
    '''

```

```

sizes = self.trainsizes
results = self.results

nmetrics = len(metrics)

# Initialize figure plots
fig, axs = plt.subplots(nmetrics, 1, figsize=(10,6))
fig.subplots_adjust(hspace=.35)
# When 1 metric is provided, allow the axs to be iterable
axs = np.array(axs).ravel()

# Construct each subplot
for metric, ax in zip(metrics, axs):
    for p, param_results in enumerate(results):
        summary = param_results['summary']
        params = param_results['params']
        # Compute the mean for multiple outputs
        res_val = np.mean(summary['val'][metric], axis=1)
        ax.plot(sizes, res_val, label=str(params))
        ax.set(ylabel=metric)
    axs[-1].set(xlabel='Train Set Size (# of folds)')
    axs[0].set(title='Validation Performance')
    axs[0].legend(bbox_to_anchor=(1.02, 1), loc='upper left',
                  ncol=1, borderaxespad=0., prop={'size': 8})
return fig, axs

def plot_best_params_by_size(self):
    '''
    Plotting function for after grid_cross_validation(), displaying
    mean (summary) train and validation set performances for the best
    parameter set for each train size for the optimized metric.

    RETURNS: the figure and axes handles
    '''
    results = self.results
    metric = self.opt_metric
    best_param_inds = self.best_param_inds
    sizes = np.array(self.trainsizes)

    # Unique set of best params for the legend
    unique_param_sets = np.unique(best_param_inds)
    lgnd_params = [self.paramsets[p] for p in unique_param_sets]

    # Initialize figure
    fig, axs = plt.subplots(2, 1, figsize=(10,6))
    fig.subplots_adjust(hspace=.35)
    # When 1 metric is provided, allow the axs to be iterable

```

```

    axs = np.array(axs).ravel()
    set_names = ['train', 'val']

    # Construct each subplot
    for i, (ax, set_name) in enumerate(zip(axs, set_names)):
        for p in unique_param_sets:
            # Obtain indices of sizes this paramset was best for
            param_size_inds = np.where(best_param_inds == p)[0]
            param_sizes = sizes[param_size_inds]
            # Compute the mean over multiple outputs for each size
            param_summary = results[p]['summary'][set_name]
            metric_scores = np.mean(param_summary[metric][param_size_inds,:
→], axis=1)

            # Plot the param results for each size it was the best for
            ax.scatter(param_sizes, metric_scores, s=120, marker=(p+2, 1))
            #ax.grid(True)

            set_name += ' Set Performance'
            ax.set(ylabel=metric, title=set_name)

        axs[-1].set(xlabel='Train Set Size (# of folds)')
        axs[0].legend(lgnd_params, bbox_to_anchor=(1.02, 1), loc='upper left',
                      ncol=1, borderaxespad=0., prop={'size': 7})
    return fig, axs

```

6 PERFORM CROSS VALIDATION

Initialize holistic cross validation objects to explore Linear, Ridge, Lasso, and ElasticNet models.

The experiments for the ElasticNet have been provided in a file (hw7_full_crossval.pkl) due to the length of time it takes to run; however, you are welcome to re-run these experiments, for all/various train set sizes, and rotations, using score_eval as the eval_func, and rmse as the metric to optimize. The file can be found in the hw7 folder in the ml_practices directory, along with this notebook.

The inputs for the models are the MI data and the outputs are the torque (you'll provide the shoulder and elbow simulataneously, as done in the previous HW).

```

[9]: """ PROVIDED
Holistic Cross Validation Options:
* ridge_alphas: list of alphas to try for the RIDGE model
* lasso_alphas: list of alphas to try for the LASSO model
* en_alphas: list of alphas to try for the ELASTICNET model
* l1_ratios: list of l1_ratios to try for the ELASTICNET model

* trainsizes: list of number of folds to utilize in the train set
* opt_metric: the optimized metric, returned by the eval_func, used

```

```

    to select the best parameter sets
* maximize_opt_metric: True if the opt_metric is maximized; False
  otherwise
* skip: the number of folds to skip when rotating through train sets
  of the same size
"""
ridge_alphas = [1, 10, 50, 100, 500, 1000, 10000]
lasso_alphas = [.001, .005, .01, .025, .05, .075, .1]
en_alphas = lasso_alphas + [0.5, 1]
l1_ratios = [0.001, .025, .05, .1, .5, 1]

trainsizes = range(1, nfolds-1)
opt_metric = 'rmse'
maximize_opt_metric = False
skip = 1

# True to always run cross validation, false to re-load existing run
# or run cross validation for the first time
force = False
# Tag for the filename to save the experiments to
prefix = "_full"

```

6.1 LINEAR REGRESSION

Ordinary least squares Linear Regression.

```

[10]: """ PROVIDED
LinearRegression

Execute cross validation procedure for all sizes for the
LinearRegression model using grid_cross_validation().
The parameter list for the LinearRegression model is a
list with just an empty dictionary [{}]
"""

lnr_fullcvfname = "hw7" + prefix + "_linear_crossval.pkl"

model = LinearRegression()
lnr_crossval = KFoldHolisticCrossValidation(model, [{}], score_eval,
                                           opt_metric, maximize_opt_metric,
                                           trainsizes, skip)

lnr_crossval_report = None
if force or (not os.path.exists(lnr_fullcvfname)):
    lnr_crossval_report = lnr_crossval.grid_cross_validation(MI_folds,
                                                            torque_folds)

    joblib.dump(lnr_crossval, lnr_fullcvfname)

```

```

else:
    # Re-load saved crossval object instead of re-running
    lnr_crossval = joblib.load(lnr_fullcvfname)
    lnr_crossval_report = {'report_by_size': lnr_crossval.report_by_size,
                          'best_param_inds': lnr_crossval.best_param_inds}

lnr_crossval.model, lnr_crossval.rotation_skip, lnr_crossval.trainsizes

```

/home/nigel/.local/lib/python3.8/site-packages/sklearn/base.py:329: UserWarning: Trying to unpickle estimator LinearRegression from version 0.23.1 when using version 0.23.2. This might lead to breaking code or invalid results. Use at your own risk.

```
warnings.warn(
```

```
[10]: (LinearRegression(), 1, range(1, 19))
```

6.2 RIDGE

$$\min_w ||y - w^T X||_2^2 + \alpha ||w||_2^2$$

α : amount of L_2 regularization to apply. Larger α greater penalize the model for larger weights

w : the weights from the model

X : feature or input data

y : true outputs

```

[11]: """ PROVIDED
      RIDGE

      Initialize a KFoldHolisticCrossValidation object that uses RIDGE
      as the model, and the provided r_allparamsets

      Execute cross validation procedure for all sizes for the Ridge
      model using grid_cross_validation()
      """

r_fullcvfname = "hw7" + prefix + "_ridge_crossval.pkl"

r_param_lists = {'alpha':ridge_alphas, 'max_iter':[1e4]}
r_allparamsets = generate_paramsets(r_param_lists)
print(pd.DataFrame(r_allparamsets))

model = Ridge()
r_crossval = KFoldHolisticCrossValidation(model, r_allparamsets, score_eval,
                                          opt_metric, maximize_opt_metric,
                                          trainsizes, skip)

```

```

r_crossval_report = None
if force or (not os.path.exists(r_fullcvfname)):
    print("Running...")
    r_crossval_report = r_crossval.grid_cross_validation(MI_folds,
                                                         torque_folds)

    joblib.dump(r_crossval, r_fullcvfname)
else:
    # Re-load saved crossval object instead of re-running
    print("Loading %s" % r_fullcvfname)
    r_crossval = joblib.load(r_fullcvfname)
    r_crossval_report = {'report_by_size' : r_crossval.report_by_size,
                        'best_param_inds': r_crossval.best_param_inds}

r_crossval.model, r_crossval.rotation_skip, r_crossval.trainsizes

```

```

alpha  max_iter
0      1      10000.0
1     10      10000.0
2     50      10000.0
3    100      10000.0
4     500      10000.0
5    1000      10000.0
6   10000      10000.0

```

Loading hw7_full_ridge_crossval.pkl

/home/nigel/.local/lib/python3.8/site-packages/sklearn/base.py:329: UserWarning:
Trying to unpickle estimator Ridge from version 0.23.1 when using version
0.23.2. This might lead to breaking code or invalid results. Use at your own
risk.

```
warnings.warn(
```

```
[11]: (Ridge(alpha=10000, max_iter=10000.0), 1, range(1, 19))
```

6.3 LASSO

$$\min_w \frac{1}{2N} \|y - w^T X\|_2^2 + \alpha \|w\|_1$$

N : the number of samples

```

[12]: """
      LASSO

      Initialize a KFoldHolisticCrossValidation object that uses LASSO
      as the model, and the provided l_allparamsets

      Execute cross validation procedure for all sizes for the Lasso
      model using grid_cross_validation()
      """

```

```

l_fullcvfname = "hw7" + prefix + "_lasso_crossval.pkl"

l_param_lists = {'alpha':lasso_alphas, 'max_iter':[1e4]}
l_allparamsets = generate_paramsets(l_param_lists)
print(pd.DataFrame(l_allparamsets))

model = Lasso()
l_crossval = KFoldHolisticCrossValidation(model, l_allparamsets, score_eval,
                                         opt_metric, maximize_opt_metric,
                                         trainsizes, skip)

l_crossval_report = None
if force or (not os.path.exists(l_fullcvfname)):
    print("Running...")
    l_crossval_report = l_crossval.grid_cross_validation(MI_folds,
                                                         torque_folds)

    # Save the cross validation object
    joblib.dump(l_crossval, l_fullcvfname)
else:
    # Re-load saved crossval object instead of re-running
    print("Loading %s" % l_fullcvfname)
    l_crossval = joblib.load(l_fullcvfname)
    l_crossval_report = {'report_by_size' : l_crossval.report_by_size,
                        'best_param_inds': l_crossval.best_param_inds}

l_crossval.model, l_crossval.rotation_skip, l_crossval.trainsizes

```

```

      alpha  max_iter
0  0.001    10000.0
1  0.005    10000.0
2  0.010    10000.0
3  0.025    10000.0
4  0.050    10000.0
5  0.075    10000.0
6  0.100    10000.0

```

Loading hw7_full_lasso_crossval.pkl

/home/nigel/.local/lib/python3.8/site-packages/sklearn/base.py:329: UserWarning:
Trying to unpickle estimator Lasso from version 0.23.1 when using version
0.23.2. This might lead to breaking code or invalid results. Use at your own
risk.

```
warnings.warn(
```

[12]: (Lasso(alpha=0.1, max_iter=10000.0), 1, range(1, 19))

6.4 ELASTICNET

$$\min_w \frac{1}{2N} \|y - w^T X\|_2^2 + \alpha L_1 \|w\|_1 + \frac{1}{2} \alpha (1 - L_1) \|w\|_2^2$$

L_1 : the L_1 ratio

```
[13]: """ PROVIDED
ELASTICNET

Initialize a KFoldHolisticCrossValidation object that uses ELASTICNET
as the model, and the provided allparamsets

Execute cross validation procedure for all sizes for the ELASTICNET
model using grid_cross_validation()

Re-load the existing experiment
"""
fullcvfname = "hw7" + prefix + "_crossval.pkl"

param_lists = {'alpha':en_alphas, 'l1_ratio':l1_ratios, 'max_iter':[1e4]}
allparamsets = generate_paramsets(param_lists)
nparamsets = len(allparamsets)
print(pd.DataFrame(allparamsets))

model = ElasticNet()
crossval = KFoldHolisticCrossValidation(model, allparamsets, score_eval,
                                       opt_metric, maximize_opt_metric,
                                       trainsizes, skip)

crossval_report = None
if force or (not os.path.exists(fullcvfname)):
    # Execute cross validation for all parameters and sizes
    print("Running...")
    crossval_report = crossval.grid_cross_validation(MI_folds,
                                                    torque_folds)

    # Save the cross validation object
    joblib.dump(crossval, fullcvfname)
else:
    print("Loading %s" % fullcvfname)
    crossval = joblib.load(fullcvfname)
    crossval_report = {'report_by_size' : crossval.report_by_size,
                      'best_param_inds': crossval.best_param_inds}

crossval.model, crossval.rotation_skip, crossval.trainsizes
```

	alpha	l1_ratio	max_iter
0	0.001	0.001	10000.0
1	0.001	0.025	10000.0
2	0.001	0.050	10000.0
3	0.001	0.100	10000.0
4	0.001	0.500	10000.0
5	0.001	1.000	10000.0

6	0.005	0.001	10000.0
7	0.005	0.025	10000.0
8	0.005	0.050	10000.0
9	0.005	0.100	10000.0
10	0.005	0.500	10000.0
11	0.005	1.000	10000.0
12	0.010	0.001	10000.0
13	0.010	0.025	10000.0
14	0.010	0.050	10000.0
15	0.010	0.100	10000.0
16	0.010	0.500	10000.0
17	0.010	1.000	10000.0
18	0.025	0.001	10000.0
19	0.025	0.025	10000.0
20	0.025	0.050	10000.0
21	0.025	0.100	10000.0
22	0.025	0.500	10000.0
23	0.025	1.000	10000.0
24	0.050	0.001	10000.0
25	0.050	0.025	10000.0
26	0.050	0.050	10000.0
27	0.050	0.100	10000.0
28	0.050	0.500	10000.0
29	0.050	1.000	10000.0
30	0.075	0.001	10000.0
31	0.075	0.025	10000.0
32	0.075	0.050	10000.0
33	0.075	0.100	10000.0
34	0.075	0.500	10000.0
35	0.075	1.000	10000.0
36	0.100	0.001	10000.0
37	0.100	0.025	10000.0
38	0.100	0.050	10000.0
39	0.100	0.100	10000.0
40	0.100	0.500	10000.0
41	0.100	1.000	10000.0
42	0.500	0.001	10000.0
43	0.500	0.025	10000.0
44	0.500	0.050	10000.0
45	0.500	0.100	10000.0
46	0.500	0.500	10000.0
47	0.500	1.000	10000.0
48	1.000	0.001	10000.0
49	1.000	0.025	10000.0
50	1.000	0.050	10000.0
51	1.000	0.100	10000.0
52	1.000	0.500	10000.0
53	1.000	1.000	10000.0

```
Loading hw7_full_crossval.pkl
```

```
/home/nigel/.local/lib/python3.8/site-packages/sklearn/base.py:329: UserWarning:  
Trying to unpickle estimator ElasticNet from version 0.23.1 when using version  
0.23.2. This might lead to breaking code or invalid results. Use at your own  
risk.
```

```
warnings.warn(  

```

```
[13]: (ElasticNet(alpha=1, l1_ratio=1, max_iter=10000.0), 1, range(1, 19))
```

7 RESULTS

7.0.1 Understand the result output structure

```
[14]: """ PROVIDED  
List KFoldHolisticCrossValidation Attributes  
"""  
dir(crossval)
```

```
[14]: ['__class__',  
      '__delattr__',  
      '__dict__',  
      '__dir__',  
      '__doc__',  
      '__eq__',  
      '__format__',  
      '__ge__',  
      '__getattr__',  
      '__gt__',  
      '__hash__',  
      '__init__',  
      '__init_subclass__',  
      '__le__',  
      '__lt__',  
      '__module__',  
      '__ne__',  
      '__new__',  
      '__reduce__',  
      '__reduce_ex__',  
      '__repr__',  
      '__setattr__',  
      '__sizeof__',  
      '__str__',  
      '__subclasshook__',  
      '__weakref__',  
      'best_param_inds',  
      'eval_func',
```

```

'get_best_params',
'get_best_params_strings',
'get_data',
'get_report_best_params_for_size',
'get_reports',
'grid_cross_validation',
'maximize_opt_metric',
'model',
'opt_metric',
'paramsets',
'perform_cross_validation',
'plot_allparams_val',
'plot_best_params_by_size',
'plot_cv',
'plot_param_train_val',
'report_by_size',
'results',
'rotation_skip',
'trainsizes']

```

```

[15]: """ PROVIDED
Results attribute is a list of dictionaries. Each element, or dictionary
corresponds to the results for a single parameter set
"""
len(crossval.results), crossval.results[0].keys()

```

```

[15]: (54, dict_keys(['params', 'results', 'summary']))

```

```

[16]: """ PROVIDED
* crossval.results[0]['results'] is a list of dictionaries with the results
  for each size for the parameter set at index 0
* crossval.results[1]['summary'] is a dictionary of summary results for the
  train, val, and test sets for the parameter set at index 1
"""
len(crossval.results[0]['results']), crossval.results[1]['summary'].keys()

```

```

[16]: (18, dict_keys(['train', 'val', 'test']))

```

```

[17]: """ PROVIDED
* crossval.results[0]['results'][2] is a dictionary with the results
  for the train size at index 2 for the parameter set at index 0
* crossval.results[1]['summary']['val'] is a dictionary of summary (over the
  sizes) results for the val set for the parameter set at index 1, for all
  metrics
"""
crossval.results[0]['results'][2].keys(), crossval.results[1]['summary']['val'].
↳keys()

```

```
[17]: (dict_keys(['train', 'val', 'test']),
      dict_keys(['mse_mean', 'mse_std', 'evar_mean', 'evar_std', 'score_mean',
                  'score_std', 'rmse_mean', 'rmse_std']))
```

```
[18]: """ PROVIDED
* crossval.results[0]['results'][2]['train'] is a dictionary of all results for
  the train set for the parameter set at index 0, the size at index 2, for all
  metrics
* crossval.results[1]['summary']['val']['mse_mean'] is a numpy array of
  → averages
  for the val set for the parameter set at index 1, for the mse. The averages
  → are
  computed over the sizes
"""
crossval.results[0]['results'][2]['train'].keys(), crossval.
→ results[1]['summary']['val']['mse_mean'].shape
```

```
[18]: (dict_keys(['mse', 'evar', 'score', 'rmse']), (18, 2))
```

```
[19]: """ PROVIDED
* crossval.results[0]['results'][2]['train']['mse'] is a dictionary of all
  results for the train set for the parameter set at index 0, the size at
  index 2, for the mse, for all rotations (there are 20 rotations when skip=1)
"""
crossval.results[0]['results'][2]['train']['mse'].shape
```

```
[19]: (20, 2)
```

7.0.2 Best Parameters for Each Size

```
[20]: """ PROVIDED
Results options:
* size_idx: index of the size from the list of train sizes to examine results
* metrics: list of summary (average) metrics to examine results
"""
# index 7 corresponds to train size 8
size_idx = 7
metrics = ['rmse_mean', 'evar_mean']
```

```
[21]: """ PROVIDED
Display the lists of the best parameter sets for each size for all
the models, expect the Linear model (as it has only one parameter set)
"""
print("Best Parameter Sets For Each Train Set Size")

print("RIDGE")
r_best_param_info = pd.DataFrame((r_crossval.trainsizes,
```

```

        r_crossval.best_param_inds,
        r_crossval.get_best_params_strings()),
        index=['train_size', 'param_index', 'paramset'])

print(r_best_param_info.T)

print("LASSO")
l_best_param_info = pd.DataFrame((l_crossval.trainsizes,
        l_crossval.best_param_inds,
        l_crossval.get_best_params_strings()),
        index=['train_size', 'param_index', 'paramset'])

print(l_best_param_info.T)

print("ELASTICNET")
best_param_info = pd.DataFrame((crossval.trainsizes,
        crossval.best_param_inds,
        crossval.get_best_params_strings()),
        index=['train_size', 'param_index', 'paramset'])

print(best_param_info.T)

```

Best Parameter Sets For Each Train Set Size

RIDGE

	train_size	param_index	paramset
0	1	5	{'alpha': 1000, 'max_iter': 10000.0}
1	2	5	{'alpha': 1000, 'max_iter': 10000.0}
2	3	5	{'alpha': 1000, 'max_iter': 10000.0}
3	4	5	{'alpha': 1000, 'max_iter': 10000.0}
4	5	5	{'alpha': 1000, 'max_iter': 10000.0}
5	6	5	{'alpha': 1000, 'max_iter': 10000.0}
6	7	5	{'alpha': 1000, 'max_iter': 10000.0}
7	8	5	{'alpha': 1000, 'max_iter': 10000.0}
8	9	5	{'alpha': 1000, 'max_iter': 10000.0}
9	10	5	{'alpha': 1000, 'max_iter': 10000.0}
10	11	5	{'alpha': 1000, 'max_iter': 10000.0}
11	12	5	{'alpha': 1000, 'max_iter': 10000.0}
12	13	5	{'alpha': 1000, 'max_iter': 10000.0}
13	14	5	{'alpha': 1000, 'max_iter': 10000.0}
14	15	5	{'alpha': 1000, 'max_iter': 10000.0}
15	16	5	{'alpha': 1000, 'max_iter': 10000.0}
16	17	5	{'alpha': 1000, 'max_iter': 10000.0}
17	18	5	{'alpha': 1000, 'max_iter': 10000.0}

LASSO

	train_size	param_index	paramset
0	1	0	{'alpha': 0.001, 'max_iter': 10000.0}
1	2	0	{'alpha': 0.001, 'max_iter': 10000.0}
2	3	0	{'alpha': 0.001, 'max_iter': 10000.0}
3	4	0	{'alpha': 0.001, 'max_iter': 10000.0}

4	5	0	{'alpha': 0.001, 'max_iter': 10000.0}
5	6	0	{'alpha': 0.001, 'max_iter': 10000.0}
6	7	0	{'alpha': 0.001, 'max_iter': 10000.0}
7	8	0	{'alpha': 0.001, 'max_iter': 10000.0}
8	9	0	{'alpha': 0.001, 'max_iter': 10000.0}
9	10	0	{'alpha': 0.001, 'max_iter': 10000.0}
10	11	0	{'alpha': 0.001, 'max_iter': 10000.0}
11	12	0	{'alpha': 0.001, 'max_iter': 10000.0}
12	13	0	{'alpha': 0.001, 'max_iter': 10000.0}
13	14	0	{'alpha': 0.001, 'max_iter': 10000.0}
14	15	0	{'alpha': 0.001, 'max_iter': 10000.0}
15	16	0	{'alpha': 0.001, 'max_iter': 10000.0}
16	17	0	{'alpha': 0.001, 'max_iter': 10000.0}
17	18	0	{'alpha': 0.001, 'max_iter': 10000.0}

ELASTICNET


	train_size	param_index	paramset
0	1	42	{'alpha': 0.5, 'l1_ratio': 0.001, 'max_iter': ...}
1	2	42	{'alpha': 0.5, 'l1_ratio': 0.001, 'max_iter': ...}
2	3	42	{'alpha': 0.5, 'l1_ratio': 0.001, 'max_iter': ...}
3	4	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
4	5	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
5	6	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
6	7	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
7	8	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
8	9	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
9	10	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
10	11	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
11	12	36	{'alpha': 0.1, 'l1_ratio': 0.001, 'max_iter': ...}
12	13	30	{'alpha': 0.075, 'l1_ratio': 0.001, 'max_iter'...
13	14	30	{'alpha': 0.075, 'l1_ratio': 0.001, 'max_iter'...
14	15	30	{'alpha': 0.075, 'l1_ratio': 0.001, 'max_iter'...
15	16	30	{'alpha': 0.075, 'l1_ratio': 0.001, 'max_iter'...
16	17	30	{'alpha': 0.075, 'l1_ratio': 0.001, 'max_iter'...
17	18	30	{'alpha': 0.075, 'l1_ratio': 0.001, 'max_iter'...

7.0.3 Plot Best Parameters for Each Size

```
[22]: """ PROVIDED
LINEAR REGRESSION
Plot the mean (summary) train and validation set performances for
each train size for the optimized metric. Use plot_best_params_by_size()


Note: for LinearRegression, there is only one parameter set.
"""
lnr_crossval.plot_best_params_by_size()
```

```
[22]: (<Figure size 720x432 with 2 Axes>,
      array([<AxesSubplot:title={'center': 'train Set Performance'},
            ylabel='rmse_mean'>,
            <AxesSubplot:title={'center': 'val Set Performance'}, xlabel='Train Set
            Size (# of folds)', ylabel='rmse_mean'>],
      dtype=object))
```



```
[23]: """ TODO
      RIDGE
      Plot the mean (summary) train and validation set performances for
      the best parameter set for each train size for the optimized
      metrics. Use plot_best_params_by_size()
      """
      r_crossval.plot_best_params_by_size()
```


```
[23]: (<Figure size 720x432 with 2 Axes>,
      array([<AxesSubplot:title={'center': 'train Set Performance'},
            ylabel='rmse_mean'>,
            <AxesSubplot:title={'center': 'val Set Performance'}, xlabel='Train Set
            Size (# of folds)', ylabel='rmse_mean'>],
      dtype=object))
```



```
[24]: """ TODO
LASSO
Plot the mean (summary) train and validation set performances for
the best parameter set for each train size for the optimized
metrics. Use plot_best_params_by_size()
"""

l_crossval.plot_best_params_by_size()
```

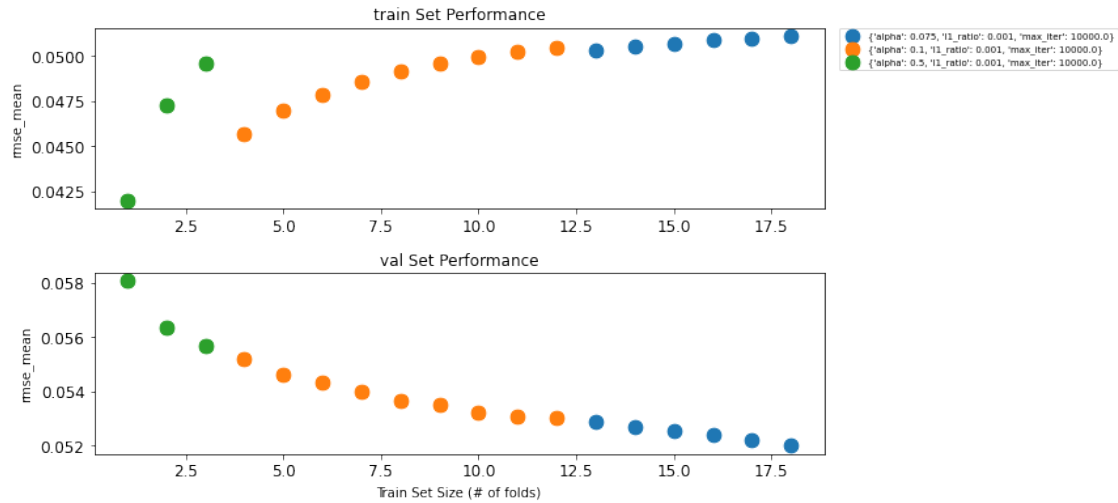
```
[24]: (<Figure size 720x432 with 2 Axes>,
array([<AxesSubplot:title={'center':'train Set Performance'},
ylabel='rmse_mean'>,
<AxesSubplot:title={'center':'val Set Performance'}, xlabel='Train Set
Size (# of folds)', ylabel='rmse_mean'>],
dtype=object))
```

```
[25]: """ TODO
ELASTICNET
Plot the mean (summary) train and validation set performances for
the best parameter set for each train size for the optimized
metrics. Use plot_best_params_by_size()
"""

crossval.plot_best_params_by_size()
```

```
[25]: (<Figure size 720x432 with 2 Axes>,
array([<AxesSubplot:title={'center':'train Set Performance'},
ylabel='rmse_mean'>,
<AxesSubplot:title={'center':'val Set Performance'}, xlabel='Train Set
Size (# of folds)', ylabel='rmse_mean'>],
dtype=object))
```




7.0.4 Plot Validation for All Parameter Sets for Each Size

```
[26]: """ TODO
      LINEAR REGRESSION
      Plot the validation results for all parameter sets over all train
      sizes, for the specified metrics, rmse_mean and evar_mean
      (this variable is declared above). Use plot_allparams_val()
      """

      lnr_crossval.plot_allparams_val(metrics)
```


```
[26]: (<Figure size 720x432 with 2 Axes>,
      array([<AxesSubplot:title={'center': 'Validation Performance'},
            ylabel='rmse_mean'>,
            <AxesSubplot:xlabel='Train Set Size (# of folds)', ylabel='evar_mean'>],
            dtype=object))
```



```
[27]: """ TODO
      RIDGE
      Plot the validation results for all parameter sets over all train
      sizes, for the specified metrics, rmse_mean and evar_mean
      (this variable is declared above). Use plot_allparams_val()
      """

      r_crossval.plot_allparams_val(metrics)
```

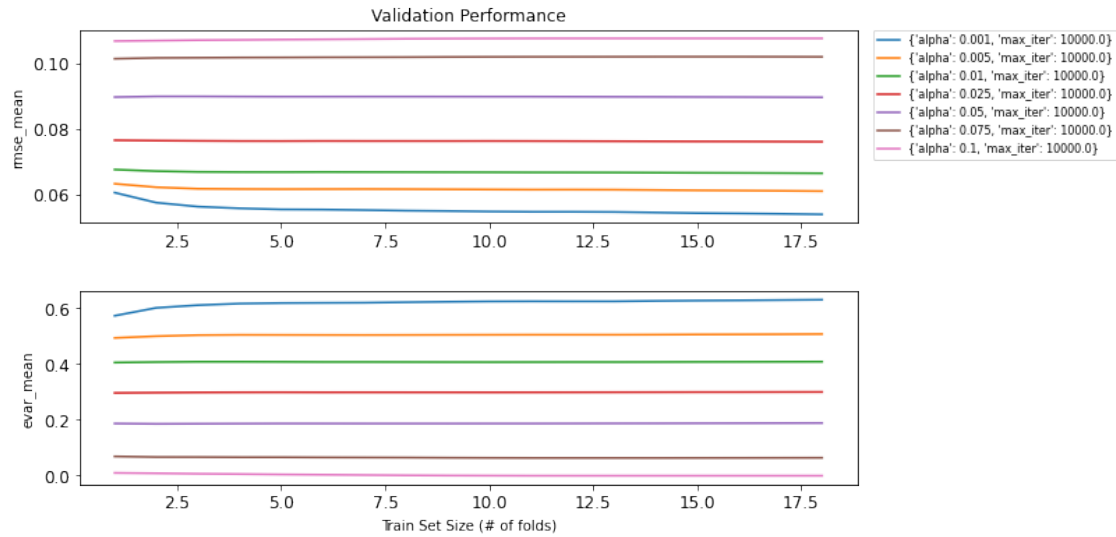
```
[27]: (<Figure size 720x432 with 2 Axes>,
      array([<AxesSubplot:title={'center':'Validation Performance'},
            ylabel='rmse_mean'>,
            <AxesSubplot:xlabel='Train Set Size (# of folds)', ylabel='evar_mean'>],
            dtype=object))
```



```
[28]: """ TODO
LASSO
Plot the validation results for all parameter sets over all train
sizes, for the specified metrics, rmse_mean and _mean_mean
(this variable is declared above). Use plot_allparams_val()
"""

l_crossval.plot_allparams_val(metrics)
```


```
[28]: (<Figure size 720x432 with 2 Axes>,
array([<AxesSubplot:title={'center':'Validation Performance'}>,
      <AxesSubplot:xlabel='Train Set Size (# of folds)', ylabel='evar_mean'>],
      dtype=object))
```



```
[29]: """ TODO
ELASTICNET
Plot the validation results for all parameter sets over all train
sizes, for the specified metrics, rmse_mean and evar_mean
(this variable is declared above). Use plot_allparams_val()
"""

crossval.plot_allparams_val(metrics)
```

```
[29]: (<Figure size 720x432 with 2 Axes>,
array([<AxesSubplot:title={'center': 'Validation Performance'}>,
      <AxesSubplot:xlabel='Train Set Size (# of folds)', ylabel='evar_mean'>],
      dtype=object))
```



7.0.5 Plot the TRAIN and VAL Set Performances

```
[30]: """ TODO
      LINEAR REGRESSION
      For the best parameter set for the train set size at
      size_idx=7 (this variable has already been declared above),
      plot the TRAIN and VAL set performances using
      plot_param_train_val() for just the optimized metric.


      Note: there is only one parameter set for the Linear model,
      thus paramidx=0
      """

      print("Train Set Size", trainsizes[size_idx])

      lnr_crossval.plot_param_train_val([lnr_crossval.opt_metric])
```

Train Set Size 8

```
[30]: (<Figure size 864x432 with 1 Axes>,
      array([<AxesSubplot:title={'center': '{}'}, xlabel='Train Set Size (# of
            folds)', ylabel='rmse_mean'>],
      dtype=object))
```




```
[31]: """ TODO
      RIDGE
      For the best parameter set for the train set size at
      size_idx=7 (this variable has already been declared above),
      plot the TRAIN and VAL set performances using
      plot_param_train_val() for just the optimized metric

      Use r_crossval.best_param_inds to get the desired parameter
      set index
      """
      print("Train Set Size", trainsizes[size_idx])
      bp_idx = r_crossval.best_param_inds[size_idx]
      r_crossval.plot_param_train_val([r_crossval.opt_metric], bp_idx)
```

Train Set Size 8

```
[31]: (<Figure size 864x432 with 1 Axes>,
      array([<AxesSubplot:title={'center': '{"alpha': 1000, 'max_iter': 10000.0}"},
            xlabel='Train Set Size (# of folds)', ylabel='rmse_mean'>],
      dtype=object))
```




```
[32]: """ TODO
LASSO
For the best parameter set for the train set size at
size_idx=7 (this variable has already been declared above),
plot the TRAIN and VAL set performances using
plot_param_train_val() for just the optimized metric
"""
print("Train Set Size", trainsizes[size_idx])

bp_idx = l_crossval.best_param_inds[size_idx]
l_crossval.plot_param_train_val([l_crossval.opt_metric], bp_idx)
```

Train Set Size 8


```
[32]: (<Figure size 864x432 with 1 Axes>,
array([<AxesSubplot:title={'center': '{alpha': 0.001, 'max_iter': 10000.0}'}],
xlabel='Train Set Size (# of folds)', ylabel='rmse_mean'>],
dtype=object))
```

```
[33]: """
      ELASTICNET
      For the best parameter set for the train set size at
      size_idx=7 (this variable has already been declared above),
      plot the TRAIN and VAL set performances using
      plot_param_train_val() for just the optimized metric
      """
      print("Train Set Size", trainsizes[size_idx])
      bp_idx = crossval.best_param_inds[size_idx]
      crossval.plot_param_train_val([crossval.opt_metric], paramidx=bp_idx)
```

Train Set Size 8

```
[33]: (<Figure size 864x432 with 1 Axes>,
      array([<AxesSubplot:title={'center': '{"alpha': 0.1, 'l1_ratio': 0.001,
'max_iter': 10000.0}"}], xlabel='Train Set Size (# of folds)',
      ylabel='rmse_mean'>],
      dtype=object))
```



7.0.6 Plot Performance over the Parameter Space

```
[34]: def plot_param_val_for_size(crossval, metric, alphas, sizeidx=0):
    ''' PROVIDED
    Plotting function for after grid_cross_validation(),
    displaying the mean (summary) train and val set performances
    for each alpha, given the size, for RIDGE and LASSO only

    PARAMS:
        crossval: cross validation object
        metric: summary metric to plot. '_mean' or '_std' must be
                append to the end of the base metric name. These
                base metric names are the keys in the dict returned
                by eval_func
        alphas: list of alpha values
        sizeidx: train size index

    RETURNS: the figure and axes handles
    '''

    sizes = crossval.trainsizes
    results = crossval.results
    best_param_inds = crossval.best_param_inds

    nalphas = len(alphas)

    nsizes = len(sizes)
    nmetrics = len(metrics)
```

```

# Initialize the matrices for the curve
Y_train = np.empty((nalphas,))
Y_val = np.empty((nalphas,))

# Obtain the mean performance for the curve
for param_res in results:
    params = param_res['params']
    summary = param_res['summary']

    alpha_idx = alphas.index(params['alpha'])

    # Compute the mean for multiple outputs
    res_train = np.mean(summary['train'][metric][sizeidx, :])
    Y_train[alpha_idx] = res_train

    res_val = np.mean(summary['val'][metric][sizeidx, :])
    Y_val[alpha_idx] = res_val

# Initialize figure plots
fig = plt.figure(figsize=(12,2))
for i, (Y, set_name) in enumerate(zip((Y_train, Y_val),
                                     ('Training', 'Validation'))):

    # Plot
    ax = fig.add_subplot(1, 2, i+1)
    ax.plot(alphas, Y)
    title = "%s Performance, Train Size %d Folds" % (set_name,
→ sizes[sizeidx])
    ax.set(title=title)
    ax.set(xlabel=r"$\alpha$", ylabel=metric)
return fig

```

```

[35]: def plot_surface(xlist, ylist, Z_train, Z_val, ylabel, zlabel,
                     elev=30, angle=45, title_suffix=""):
    ''' PROVIDED
    Helper plotting function. x-axis is always alpha

    REQUIRES: from mpl_toolkits.mplot3d import Axes3D

    PARAMS:
        xlist: list of x values
        ylist: list of y values
        Z_train: matrix of performance results from the training set
        Z_val: matrix of performance results from the validation set
        ylabel: y-axis label
        zlabel: z-axis label
        elev: elevation of the 3D plot for the view

```

```

        angle: angle in degrees of the 3D plot for the view
        title_suffix: string to append to each subplot title

    RETURNS: the figure and axes handles
    '''
    # Initialize figure
    fig = plt.figure(figsize=(15,5))
    X, Y = np.meshgrid(xlist, ylist)
    for i, (Z, set_name) in enumerate(zip((Z_train, Z_val),
                                         ('Training', 'Validation'))):

        # Plot the surface
        ax = fig.add_subplot(1, 2, i+1, projection='3d')
        surf = ax.plot_surface(X, Y, Z, cmap=cm.coolwarm,
                               linewidth=0, antialiased=False)
        title = "%s Performance %s" % (set_name, title_suffix)
        ax.view_init(elev=elev, azimuth=angle)
        ax.set(title=title)
        ax.set(xlabel=r"$\alpha$", ylabel=ylabel, zlabel=zlabel)
    return fig

```

```

[36]: def plot_param_val_surface_RL(crossval, metric, alphas, elev=30, angle=245):
    ''' PROVIDED
    Plotting function for after grid_cross_validation(),
    displaying the mean (summary) train and val set performances
    for each alpha, for all sizes, for RIDGE and LASSO only

    REQUIRES: from mpl_toolkits.mplot3d import Axes3D

    PARAMS:
        crossval: cross validation object
        metric: summary metric to plot. '_mean' or '_std' must be
                append to the end of the base metric name. These
                base metric names are the keys in the dict returned
                by eval_func
        alphas: list of alpha values
        elev: elevation of the 3D plot for the view
        angle: angle in degrees of the 3D plot for the view

    RETURNS: the figure and axes handles
    '''
    sizes = crossval.trainsizes
    results = crossval.results
    best_param_inds = crossval.best_param_inds

    nalphas = len(alphas)

    nsizes = len(sizes)

```

```

nmetrics = len(metrics)

# Initialize the matrices for the surface
Z_train = np.empty((nsizes, nalphas))
Z_val = np.empty((nsizes, nalphas))

# Obtain the mean performance for the surface
for param_res in results:
    params = param_res['params']
    summary = param_res['summary']

    alpha_idx = alphas.index(params['alpha'])

    # Compute the mean for multiple outputs
    res_train = np.mean(summary['train'][metric], axis=1)
    Z_train[:, alpha_idx] = res_train

    # Compute the mean for multiple outputs
    res_val = np.mean(summary['val'][metric], axis=1)
    Z_val[:, alpha_idx] = res_val

fig = plot_surface(alphas, sizes, Z_train, Z_val, 'size (# of folds)',
                  metric, elev, angle)

return fig

```

```

[37]: def plot_param_val_surface_EN(crossval, metric, param_lists,
                                   sizeidx=0, elev=35, angle=280):

    ''' PROVIDED
    Plotting function for after grid_cross_validation(),
    displaying the mean (summary) train and val set performances
    for each alpha and l1_ratio, given the size, for the ELASTICNET

    REQUIRES: from mpl_toolkits.mplot3d import Axes3D

    PARAMS:
        crossval: cross validation object
        metric: summary metric to plot. '_mean' or '_std' must be
                append to the end of the base metric name. These
                base metric names are the keys in the dict returned
                by eval_func
        param_lists: dictionary of the list of alphas and l1_ratios
        sizeidx: train size index
        elev: elevation of the 3D plot for the view
        angle: angle in degrees of the 3D plot for the view

    RETURNS: the figure and axes handles
    '''

```

```

sizes = crossval.trainsizes
results = crossval.results
best_param_inds = crossval.best_param_inds

alphas = list(param_lists['alpha'])
l1_ratios = list(param_lists['l1_ratio'])

nalphas = len(alphas)
nl1_ratios = len(l1_ratios)

nsizes = len(sizes)
nmetrics = len(metrics)

# Initialize the matrices for the surface
Z_train = np.empty((nl1_ratios, nalphas))
Z_val = np.empty((nl1_ratios, nalphas))

# Obtain the mean performance for the surface
for param_res in results:
    params = param_res['params']
    summary = param_res['summary']

    alpha_idx = alphas.index(params['alpha'])
    l1_idx = l1_ratios.index(params['l1_ratio'])

    # Compute the mean for multiple outputs
    res_train = np.mean(summary['train'][metric][sizeidx, :])
    Z_train[l1_idx, alpha_idx] = res_train

    res_val = np.mean(summary['val'][metric][sizeidx, :])
    Z_val[l1_idx, alpha_idx] = res_val

fig = plot_surface(alphas, l1_ratios, Z_train, Z_val, 'l1_ratio',
                  metric, elev, angle, ', Size %d Folds' % sizes[sizeidx])

return fig

```

```

[38]: """ PROVIDED
List the parameter sets explored for RIDGE
"""
r_crossval.paramsets

```

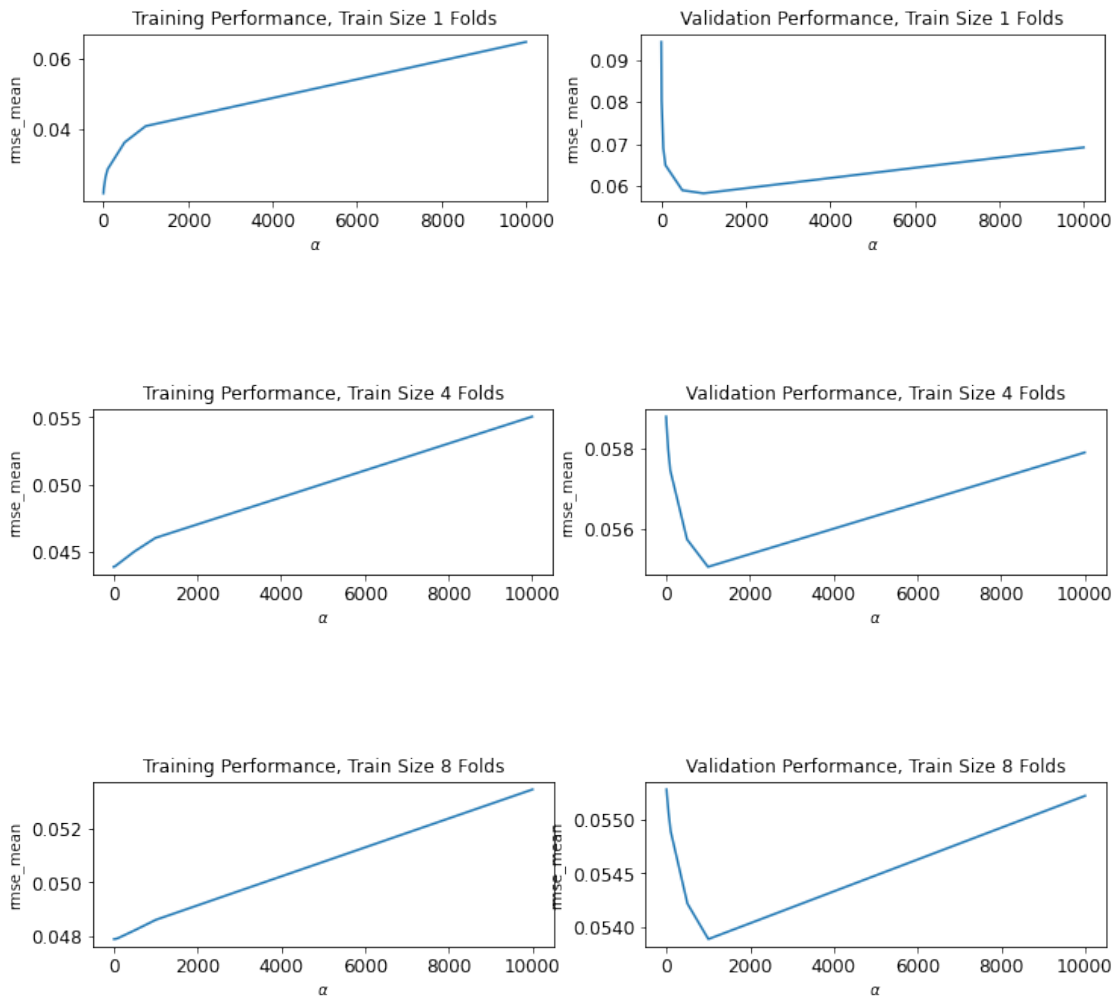
```

[38]: [{'alpha': 1, 'max_iter': 10000.0},
{'alpha': 10, 'max_iter': 10000.0},
{'alpha': 50, 'max_iter': 10000.0},
{'alpha': 100, 'max_iter': 10000.0},
{'alpha': 500, 'max_iter': 10000.0},
{'alpha': 1000, 'max_iter': 10000.0},

```

```
{'alpha': 10000, 'max_iter': 10000.0}]
```

```
[71]: """ TODO
Plot the performance versus alpha for the RIDGE model
using plot_param_val_for_size() for size indices 0, 3, and 7,
for the optimized metric (use r_crossval.opt_metric)
"""
size_indices = [0, 3, size_idx]
for si in size_indices:
    plot_param_val_for_size(r_crossval, r_crossval.opt_metric,
        ↪r_param_lists['alpha'], si)
```



```
[59]: """ TODO
RIDGE
Use plot_param_val_surface_RL() to plot the surface of the training
and validation set performance versus alpha and size in the X and Y axes,
```

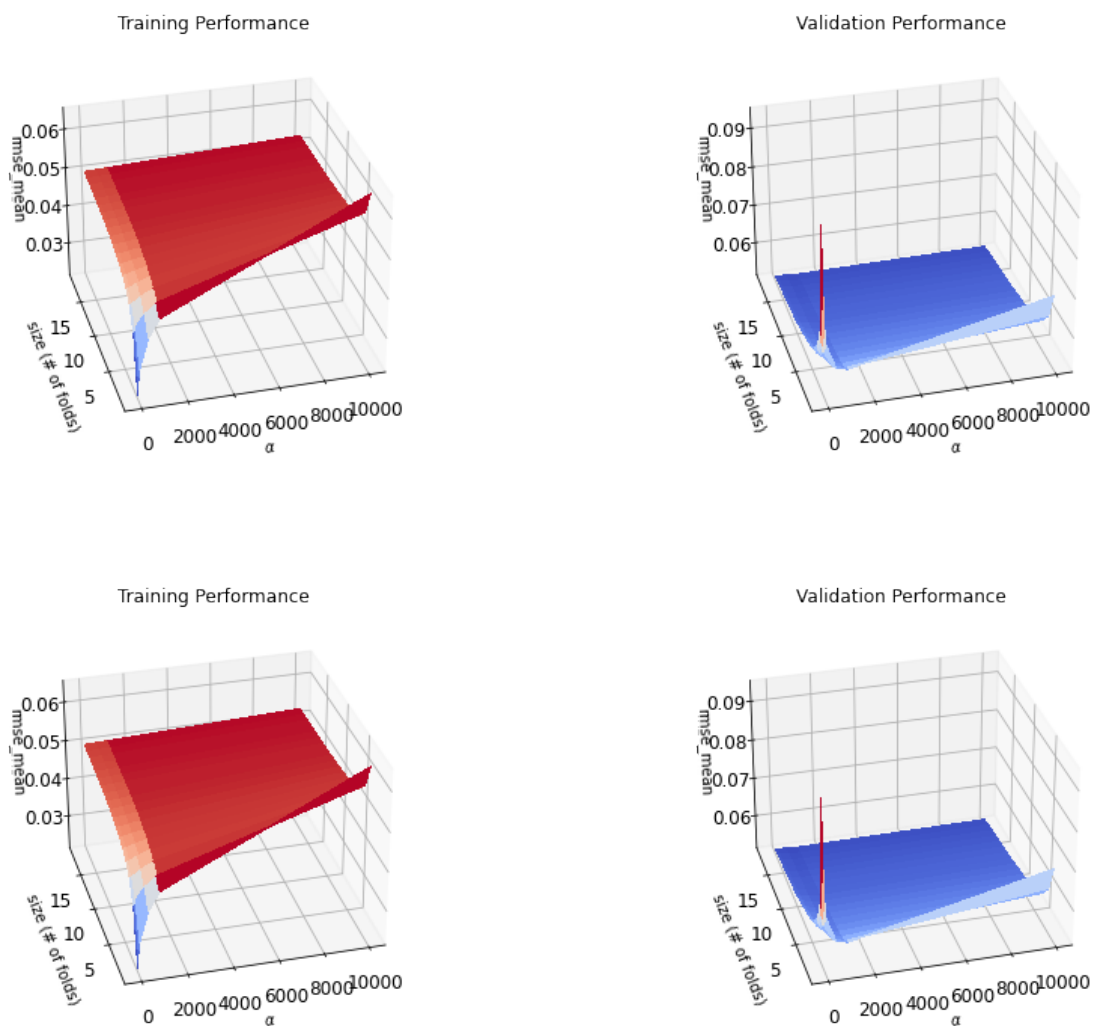
```

using the optimized metric
"""
# Feel free to adjust these to understand the shape of the surface
# Elevation of the plot
elev = 30
# Angle the plot is viewed
angle = 255

# TODO: Plot
plot_param_val_surface_RL(r_crossval, r_crossval.opt_metric,
    ↪r_param_lists['alpha'], elev, angle)

```

[59]:



```

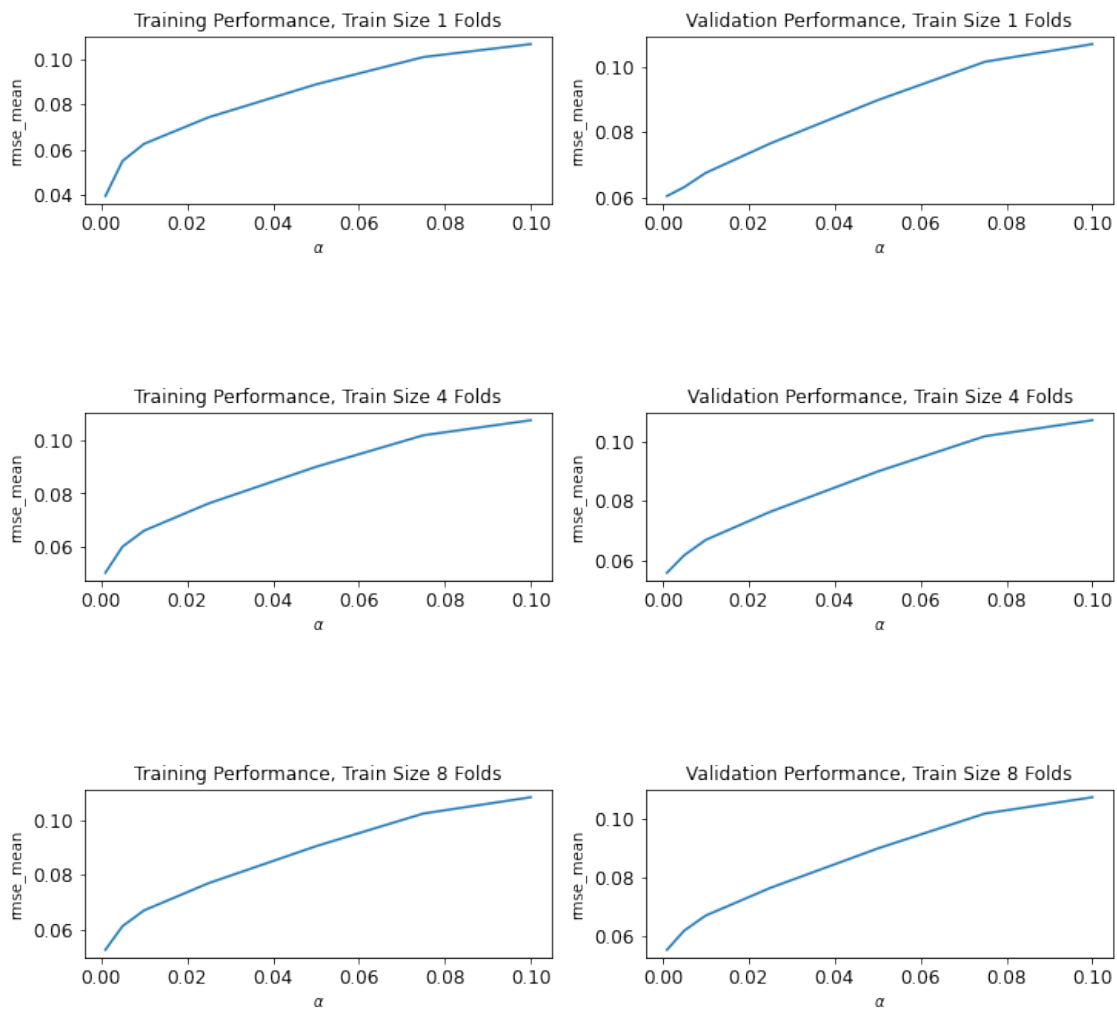
[ ]: """ PROVIDED
List the parameter sets explored for LASSO
"""

```



```
l_crossval.paramsets
```

```
[72]: """ TODO
Plot the performance versus alpha for the LASSO model
using plot_param_val_for_size() for size indices 0, 3, and 7,
for the optimized metric
"""
size_indices = [0, 3, size_idx]
for si in size_indices:
    plot_param_val_for_size(l_crossval, l_crossval.opt_metric,
        ↪l_param_lists['alpha'], si)
```



```
[68]: """ TODO
LASSO
Use plot_param_val_surface_RL() to plot the surface of the training
and validation set performance versus alpha and size in the X and Y axes,
```

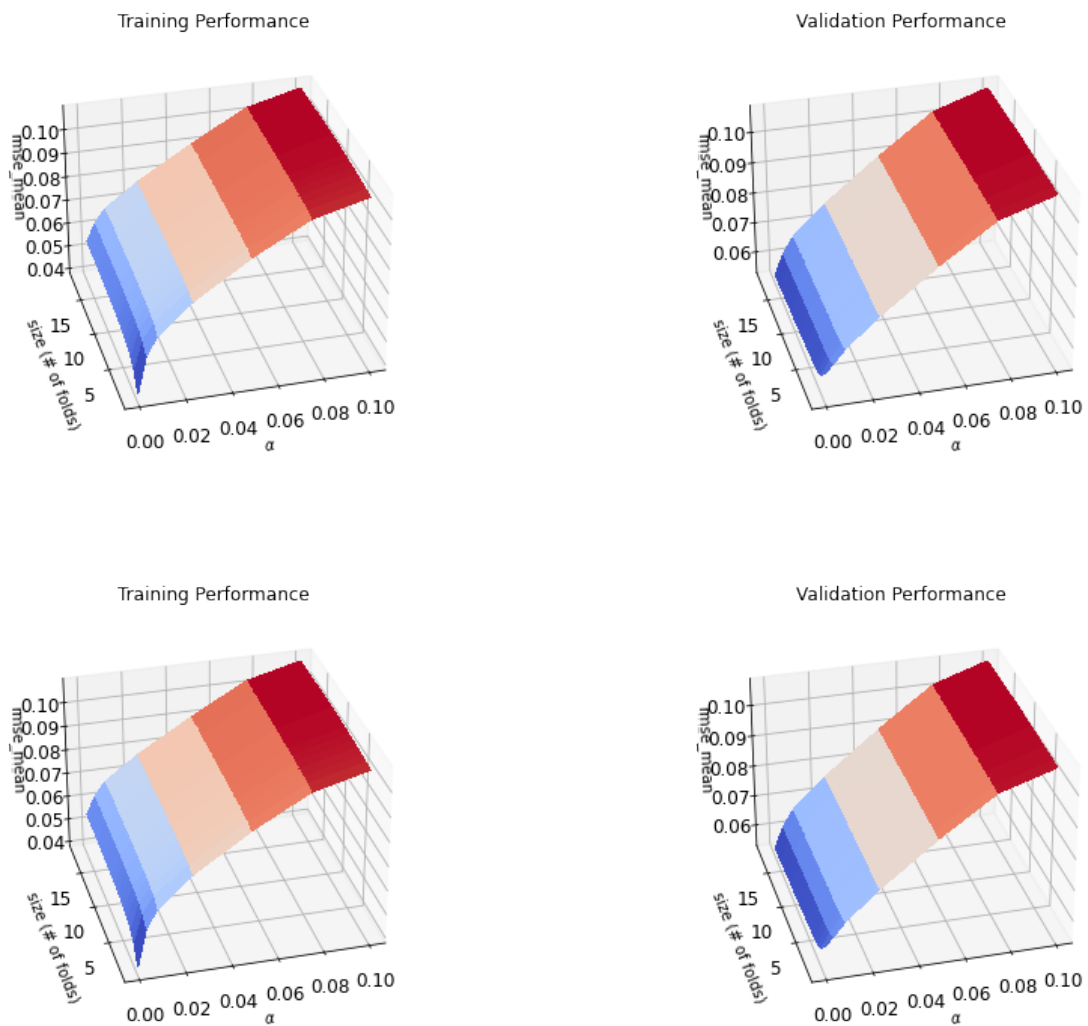
```

using the optimized metric
"""
# Feel free to adjust these to understand the shape of the surface
# Elevation of the plot
elev = 30
# Angle the plot is viewed
angle = 255

# TODO: Plot
plot_param_val_surface_RL(l_crossval, l_crossval.opt_metric,
    ↪l_param_lists['alpha'], elev, angle)

```

[68]:



```

[ ]: """ PROVIDED
List the parameter sets explored for ELASTICNET
"""

```


```
crossval.paramsets
```

```
[73]: """ TODO
ELASTICNET
Use plot_param_val_surface_EN() to plot the surface of the training
and validation set performance versus alpha and l1_ratio in the X
and Y axes for the size indices of 0, 3, and 7, for crossval.opt_metric
"""


# Feel free to adjust these to understand the shape of the surface
# Elevation of the plot
elev = 25
# Angle the plot is viewed
angle = 280

# TODO: Plot
size_indices = [0, 3, size_idx]
for si in size_indices:
    plot_param_val_surface_EN(crossval, crossval.opt_metric, param_lists,
                              si, elev, angle)
```

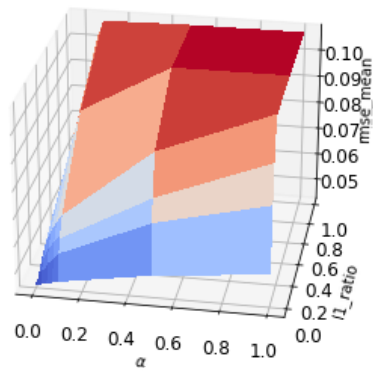
Training Performance , Size 1 Folds




Validation Performance , Size 1 Folds




Training Performance , Size 4 Folds




Validation Performance , Size 4 Folds



Training Performance , Size 8 Folds



Validation Performance , Size 8 Folds



7.0.7 Paired t-tests

We can use paired t-tests to assess statistical significant differences between the mean test set performances of the models

```
[70]: """ PROVIDED
Obtain all the results for all the models
"""

# LinearRegression
lnr_all_results = lnr_crossval.results

# RIDGE
r_all_results = r_crossval.results

# LASSO
l_all_results = l_crossval.results
```

```
# ELASTICNET
all_results = crossval.results
```

```
[74]: """ TODO
Complete the plotting code

Plot distributions of the Validation and Test scores from the
best parameter set for each base model for the corresponding
size indices, [0, 3, 7]. The metric of interest is rmse.
These are the distribution of results from each rotation of
the training set
"""

metric = 'rmse'
set_names = ['val', 'test']
nbins = 11

# Size indices
size_indices = [0, 3, size_idx]

for si in size_indices:
    # Obtain the index of the best parameter set for the size
    # RIDGE
    r_bp_idx = r_crossval.best_param_inds[si]
    # LASSO
    l_bp_idx = l_crossval.best_param_inds[si]
    # ELASTICNET
    bp_idx = crossval.best_param_inds[si]

    # Construct the figure
    fig, axs = plt.subplots(2, 2, figsize=(15,7))
    for i, set_name in enumerate(set_names):
        title = '%s, Size %d' % (set_name, trainsizes[si])

        # LINEAR
        # Note: there's only 1 parameter set for the Linear model
        lnr_res = lnr_all_results[0]['results'][si][set_name]
        lnr_scores = np.mean(lnr_res[metric], axis=1)

        # RIDGE
        # Obtain results for the best parameter set for the size
        ridge_res = r_all_results[r_bp_idx]['results'][si][set_name]
        # Compute the mean of the outputs for each data set rotation
        ridge_scores = np.mean(ridge_res[metric], axis=1)

        # LASSO
        lasso_res = l_all_results[l_bp_idx]['results'][si][set_name]
```

```

lasso_scores = np.mean(lasso_res[metric], axis=1)

# ELASTICNET
res = all_results[bp_idx]['results'][si][set_name]
elastic_scores = np.mean(res[metric], axis=1)

# Determine the edges for the bins in the histograms
all_scores = np.concatenate((elastic_scores, ridge_scores,
                              lasso_scores, lnr_scores))

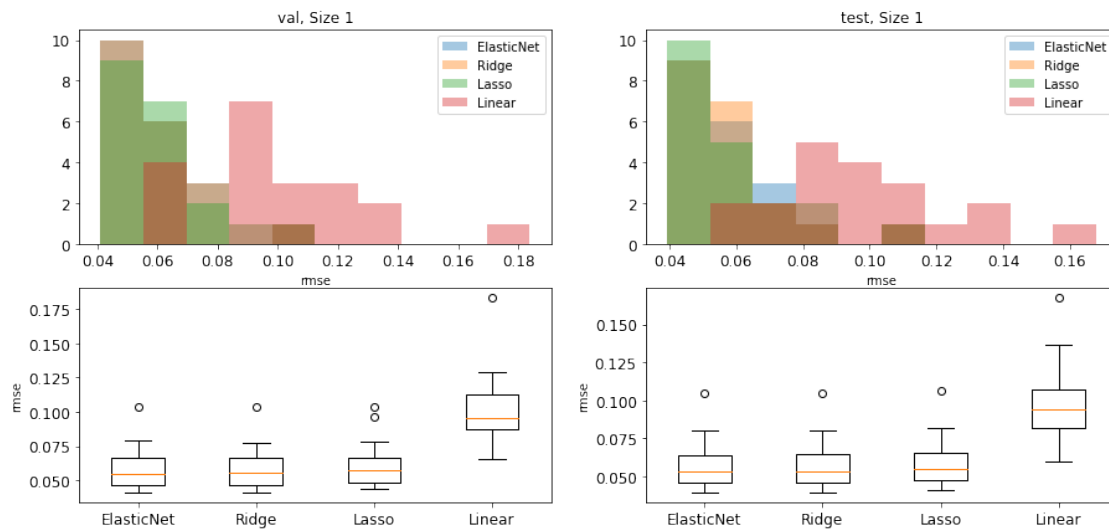
mn = np.min(all_scores)
mx = np.max(all_scores)
bins = np.linspace(mn, mx, nbins)

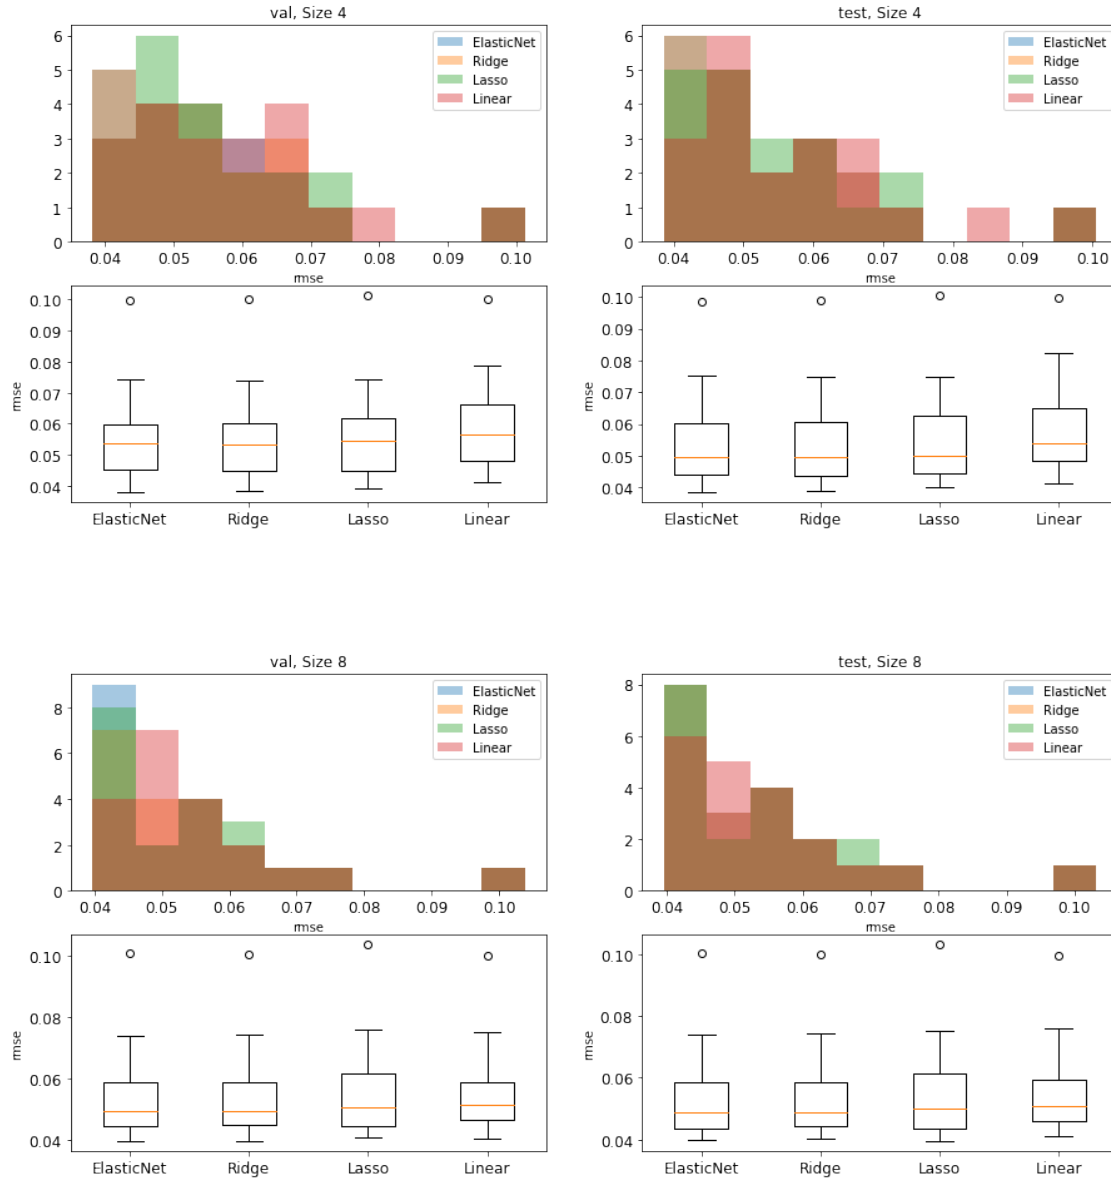
# Histograms
# TODO: include the hist of the elastic net scores

axs[0, i].hist(elastic_scores, bins=bins, alpha=.4)
axs[0, i].hist(ridge_scores, bins=bins, alpha=.4)
axs[0, i].hist(lasso_scores, bins=bins, alpha=.4)
axs[0, i].hist(lnr_scores, bins=bins, alpha=.4)
axs[0, i].legend(['ElasticNet', 'Ridge', 'Lasso', 'Linear'])
axs[0, i].set(title=title, xlabel=metric)

# Boxplots
axs[1, i].boxplot([elastic_scores, ridge_scores, lasso_scores,
↪lnr_scores])
axs[1, i].set_xticklabels(['ElasticNet', 'Ridge', 'Lasso', 'Linear'])
axs[1, i].set(ylabel=metric)

```





```
[75]: """ TODO
Dependent Sample Paired t-test
Two-sided t-test for the null hypothesis that mean of the distribution
of differences between the two test performance distributions is zero
"""
print("Train Set Size", trainsizes[size_idx])

# LINEAR
# Note: there's only 1 parameter set for the LinearRegression model
lnr_res = lnr_crossval.results[0]['results'][size_idx]['test']
lnr_test_res = np.mean(lnr_res[metric], axis=1)
```

```

# RIDGE
# Obtain index of best parameters for train size 8
r_bp_idx = r_crossval.best_param_inds[size_idx]
# Obtain all results for the best parameter set for train size 8
ridge_res = r_all_results[r_bp_idx]['results'][size_idx]['test']
# Compute the mean of the outputs for each data set rotation
ridge_test_res = np.mean(ridge_res[metric], axis=1)

# LASSO
l_bp_idx = l_crossval.best_param_inds[size_idx]
lasso_res = l_all_results[l_bp_idx]['results'][size_idx]['test']
lasso_test_res = np.mean(lasso_res[metric], axis=1)

# TODO: ELASTICNET
bp_idx = crossval.best_param_inds[size_idx]
net_res = all_results[bp_idx]['results'][size_idx]['test']
elastic_test_res = np.mean(net_res[metric], axis=1)

```

Train Set Size 8

```

[93]: """ TODO
      ELASTICNET vs RIDGE
      Execute the paired t-test to determine whether to reject the null hypothesis
      (i.e.  $H_0$ ) with 95% confidence.  $H_0$  is that the mean of the distribution of the
      differences between test scores for the best ELASTICNET model and the best
      ↪ RIDGE
      is zero, when using a training size of 8 (i.e. the size at index 7 of the
      trainsizes list). Display the t-statistic, the p-value, and the mean of the
      differences (i.e.  $\text{mean}(\text{elastic\_test\_res} - \text{ridge\_test\_res})$ )

      Use stats.ttest_rel(). See the API reference above.
      Do the same for all the pairing of models
      """
      t_statistic, p_value = stats.ttest_rel(elastic_test_res, ridge_test_res)
      print('t-statistic', t_statistic)
      print('p-value', p_value)
      print('mean', np.mean(elastic_test_res - ridge_test_res))

```

```

t-statistic -2.4787147850144167
p-value 0.022736733532476117
mean -0.00021785532084971117

```

[]:

```

[94]: """ TODO
      ELASTICNET vs LASSO

```



```
Execute the paired t-test
"""
```

```
t_statistic, p_value = stats.ttest_rel(elastic_test_res, lasso_test_res)
print('t-statistic', t_statistic)
print('p-value', p_value)
print('mean', np.mean(elastic_test_res - lasso_test_res))
```

```
t-statistic -4.3151954381636095
p-value 0.00037323234323962415
mean -0.0014263082520257705
```

```
[ ]:
```

```
[95]: """ TODO
ELASTICNET vs LinearRegression
Execute the paired t-test
"""
```

```
t_statistic, p_value = stats.ttest_rel(elastic_test_res, lnr_test_res)
print('t-statistic', t_statistic)
print('p-value', p_value)
print('mean', np.mean(elastic_test_res - lnr_test_res))
```

```
t-statistic -4.873596012694302
p-value 0.00010549817316970218
mean -0.0015881764958851166
```

```
[ ]:
```

```
[96]: """ TODO
RIDGE vs LASSO
Execute the paired t-test
"""
```

```
t_statistic, p_value = stats.ttest_rel(ridge_test_res, lasso_test_res)
print('t-statistic', t_statistic)
print('p-value', p_value)
print('mean', np.mean(ridge_test_res - lasso_test_res))
```

```
t-statistic -2.9120174293543895
p-value 0.008939372406479561
mean -0.0012084529311760593
```

```
[ ]:
```

```
[97]: """ TODO
RIDGE vs LinearRegression
Execute the paired t-test
"""
```

```
t_statistic, p_value = stats.ttest_rel(ridge_test_res, lnr_test_res)
print('t-statistic', t_statistic)
print('p-value', p_value)
print('mean', np.mean(ridge_test_res - lnr_test_res))
```

```
t-statistic -5.610555129622476
p-value 2.071569385430273e-05
mean -0.0013703211750354053
```

[]:

```
[98]: """ TODO
      LASSO vs LinearRegression
      Execute the paired t-test
      """

      t_statistic, p_value = stats.ttest_rel(lasso_test_res, lnr_test_res)
      print('t-statistic', t_statistic)
      print('p-value', p_value)
      print('mean', np.mean(lasso_test_res - lnr_test_res))
```

```
t-statistic -0.2528166363086273
p-value 0.8031249544779376
mean -0.00016186824385934598
```

[]:

8 DISCUSSION

For each question write one brief paragraph of discussion:

1. Interpret the meaning of the t-test results using 95% confidence. Discuss the statistical meaning as well as the practical interpretation of the results in the context of the data set.
2. For the Elastic Net Model, discuss the differences in the surfaces between the train sizes of 1, 4, and 8 folds, for both the training and validation sets.
3. For each of the train set sizes of 1, 4, and 8 folds, which model (Linear, Lasso, Ridge, or ElasticNet) and corresponding parameter set would you select and why? Specify which model and parameter set for each size. For each size, use `plot_param_train_val()` to view the train, val, and test sets of the chosen model(s). Remember, selections should be made based on the validation performance.


9 Answers

1. If the p-value is greater greater than 0.05 then we cannot reject the null hypothesis. If the p-value is smaller than 5%, then we reject the null hypothesis. With this we see that ELASTICNET vs RIDGE, ELASTICNET vs LASSO, ELASTICNET vs LinearRegression, and RIDGE vs LASSO all have p-values that are less than 0.05.

2. For validation set, the surface is different for each fold at alpha 0.0 and l1 ratios. For for training, folds 8 and 4 have different l1 ratios than fold 1. 8 and 4 have a surface between .4 and .6 while 1 fold doesn't
3. I would pick elasticnet with best paramset because it is closest to the test data.

```
[101]: bp_idx = crossval.best_param_inds[size_idx]
crossval.plot_param_train_val([crossval.opt_metric], paramidx=bp_idx,
↪view_test=True)
```

```
[101]: (<Figure size 864x432 with 1 Axes>,
array([<AxesSubplot:title={'center': '{alpha': 0.1, 'l1_ratio': 0.001,
'max_iter': 10000.0}"}], xlabel='Train Set Size (# of folds)',
ylabel='rmse_mean'>],
dtype=object))
```



```
[ ]:
```